

Machine-Level Programming II: Procedures and Data

Computer Systems
Friday, October 13, 2023

Today

Procedures

- **Stack Structure**
- **Calling Conventions**
 - Passing control
 - Passing data
 - Managing local data
- **Illustration of Recursion**

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

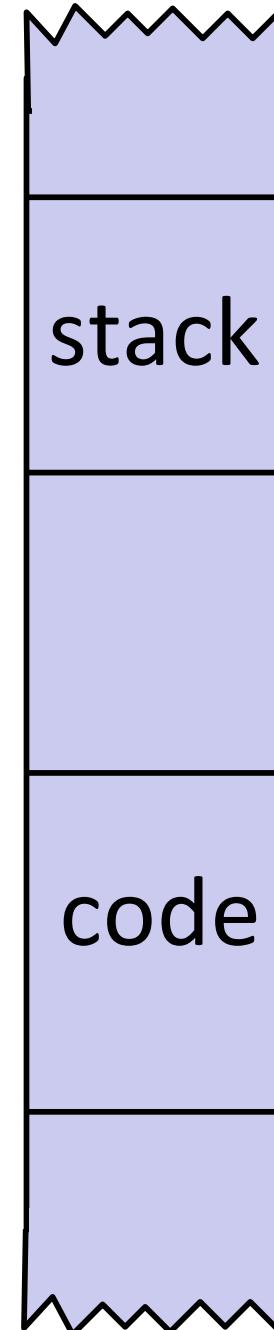
Structures

- Allocation
- Access
- Alignment

x86-64 Stack

**Region of memory managed
with stack discipline**

- Memory viewed as array of bytes.
- Different regions have different purposes.
- (Like ABI, a policy decision)



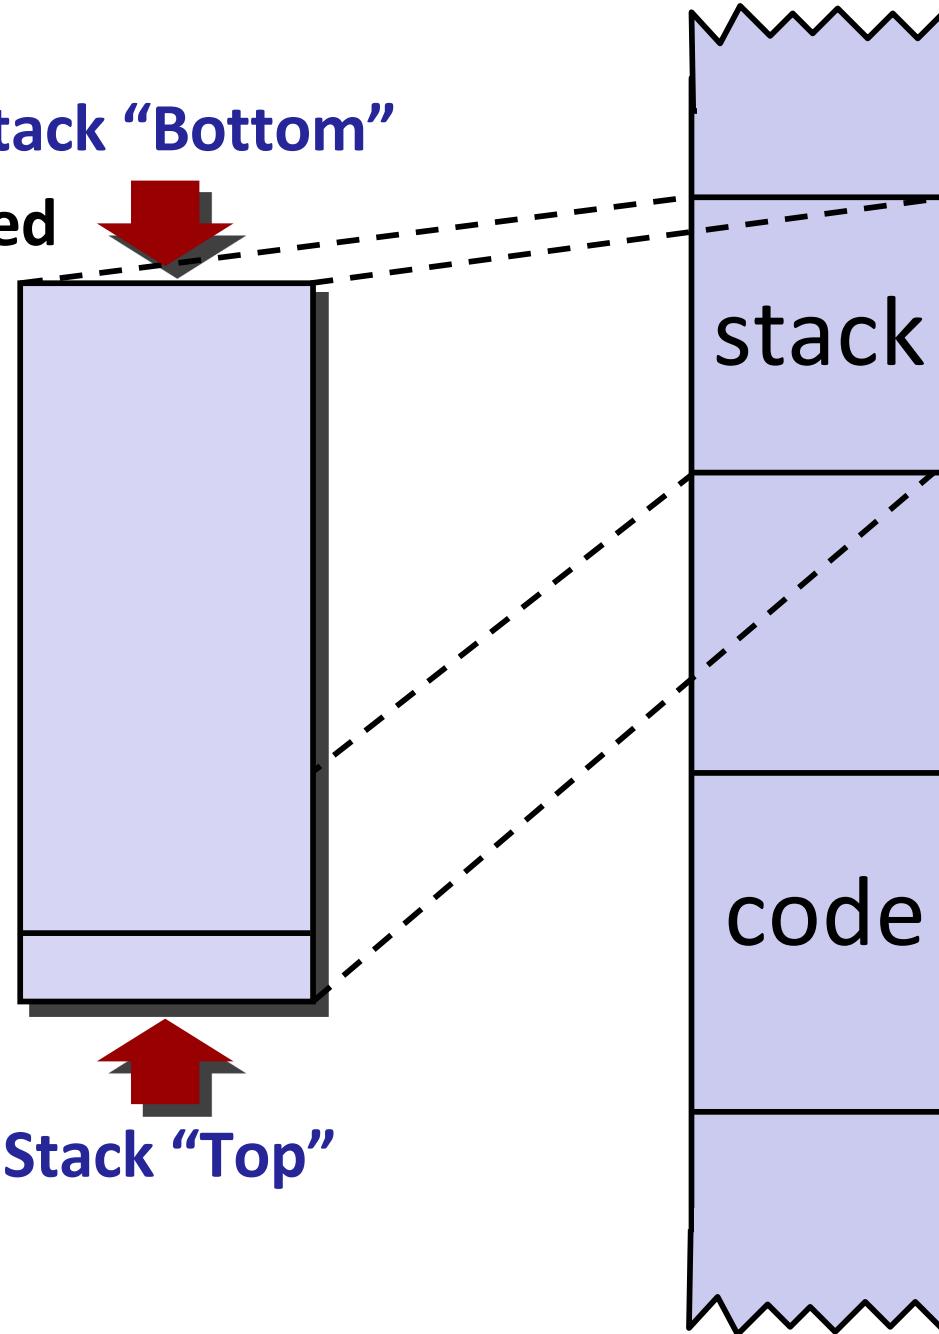
x86-64 Stack

Region of memory managed
with stack discipline

Stack “Bottom”

Stack Pointer: %rsp →

Stack “Top”



x86-64 Stack

**Region of memory managed
with stack discipline**

Grows toward lower addresses

**Register `%rsp` contains
lowest stack address**

- address of “top” element

Stack Pointer: `%rsp` →

Stack “Bottom”



Stack “Top”

x86-64 Stack: Push

pushq Src

- Fetch operand at *Src*
- Decrement **%rsp** by 8
- Write operand at address given by **%rsp**



Stack Pointer: →
%rsp

Stack “Bottom”



Stack “Top”

x86-64 Stack: Push

pushq Src

- Fetch operand at *Src*
- Decrement **%rsp** by 8
- Write operand at address given by **%rsp**

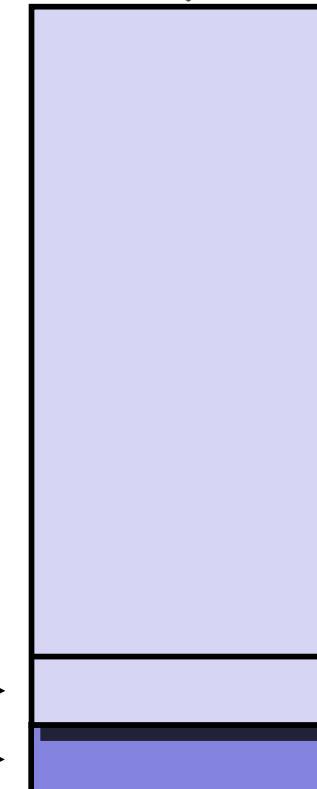
val

Stack Pointer:

%rsp

↓ -8 →

Stack “Bottom”



Increasing Addresses
↑

Stack Grows Down
↓

Stack “Top”

x86-64 Stack: Pop

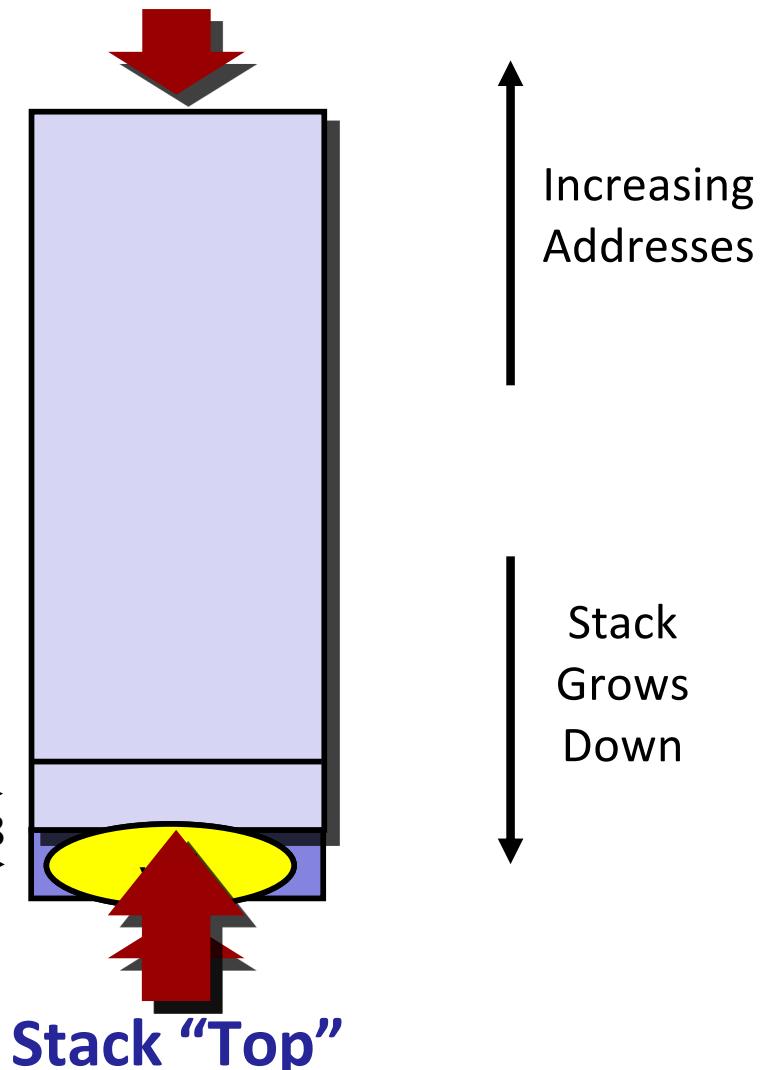
popq Dest

- Read value at address given by `%rsp`
- Increment `%rsp` by 8
- Store value at Dest (usually a register)

Value is **copied**; it remains
in memory at old `%rsp`

Stack Pointer:
`%rsp` 

Stack “Bottom”



Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- Illustration of Recursion

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

Code Examples

```
void multstore(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

000000000400540 <multstore>:

```
400540: push    %rbx          # Save %rbx
400541: mov     %rdx,%rbx    # Save dest
400544: call    400550 <mult2> # mult2(x,y)
400549: mov     %rax,(%rbx)   # Save at dest
40054c: pop    %rbx          # Restore %rbx
40054d: ret             # Return
```

```
long mult2(long a, long b)
{
    long s = a * b;
    return s;
}
```

000000000400550 <mult2>:

```
400550: mov     %rdi,%rax    # a
400553: imul   %rsi,%rax    # a * b
400557: ret             # Return
```

Procedure Control Flow

Use stack to support procedure call and return

Procedure call: `call label`

- Push return address on stack
- Jump to *label*

Return address:

- Address of the next instruction right after call
- Example from disassembly

Procedure return: `ret`

- Pop address from stack
- Jump to address

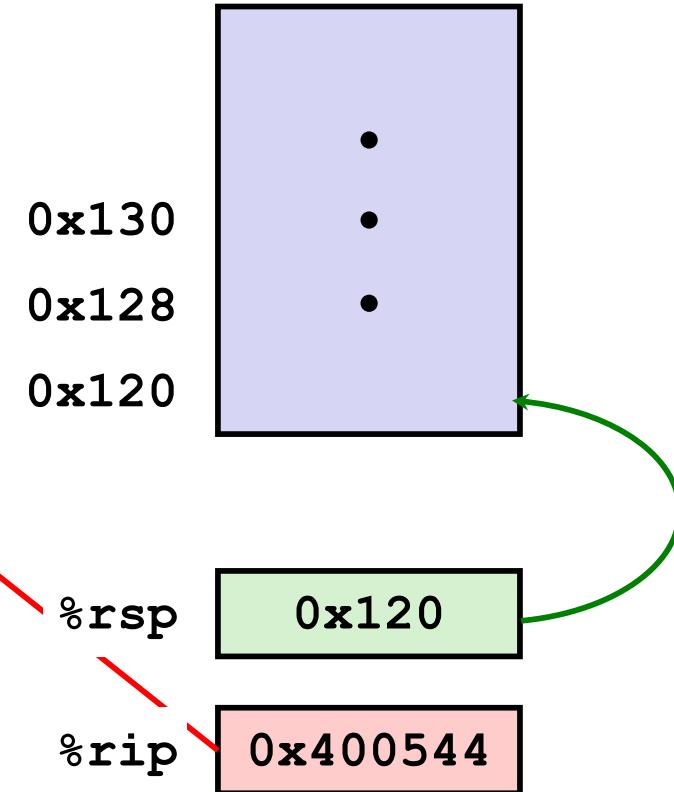
These instructions are sometimes printed with a q suffix

- This is just to remind you that you're looking at 64-bit code

Control Flow Example #1

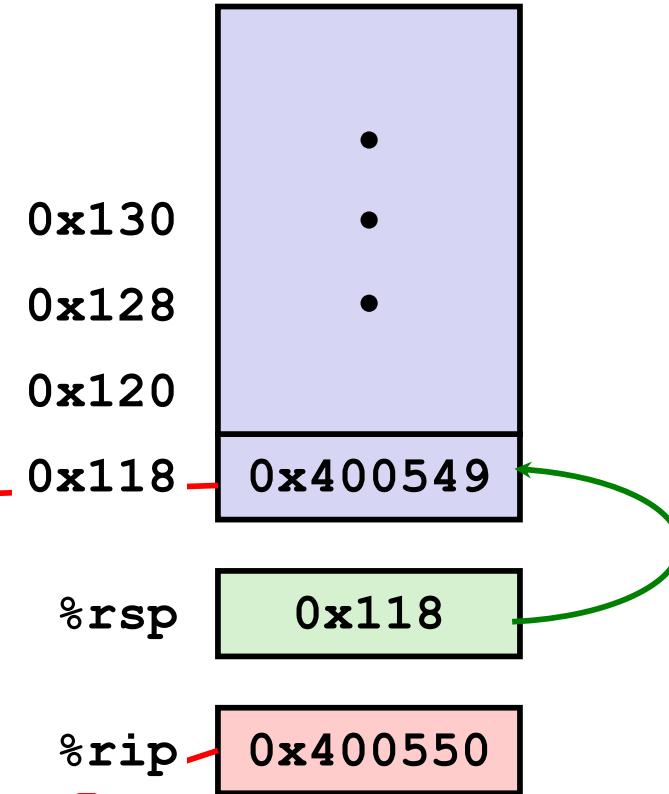
```
0000000000400540 <multstore>:  
.  
. .  
400544: call    400550 <mult2>  
400549: mov     %rax, (%rbx)  
. .
```

```
0000000000400550 <mult2>:  
400550: mov     %rdi,%rax  
. .  
400557: ret
```



Control Flow Example #2

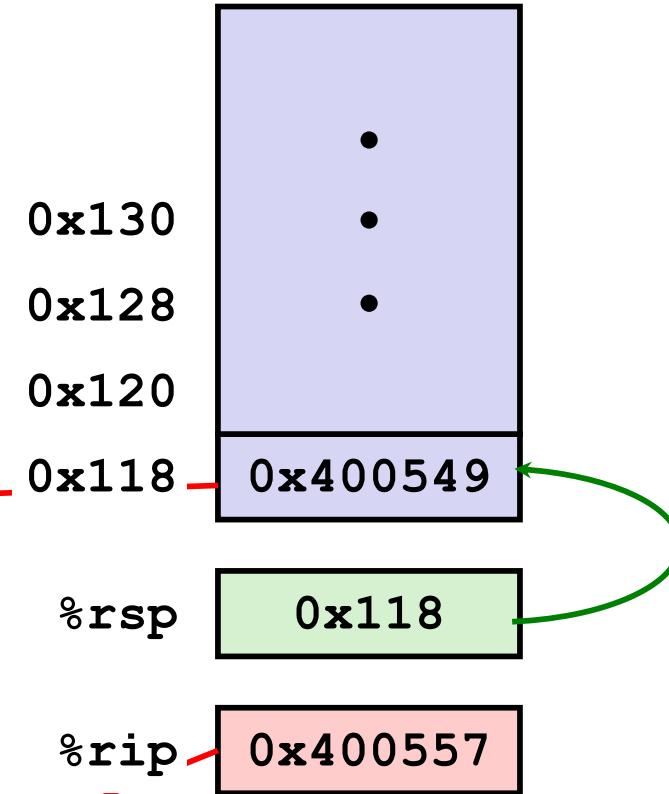
```
0000000000400540 <multstore>:  
.  
. .  
400544: call    400550 <mult2>  
400549: mov     %rax, (%rbx) ←
```



```
0000000000400550 <mult2>:  
400550: mov     %rdi,%rax ←  
. .  
400557: ret
```

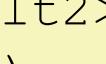
Control Flow Example #3

```
0000000000400540 <multstore>:  
.  
. .  
400544: call    400550 <mult2>  
400549: mov     %rax, (%rbx) ←
```

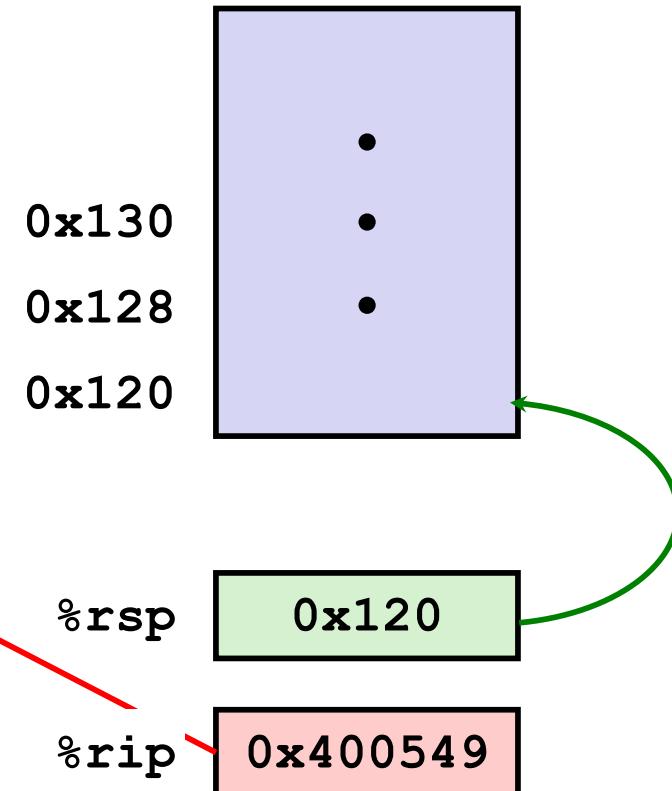


```
0000000000400550 <mult2>:  
400550: mov     %rdi,%rax  
. .  
400557: ret ←
```

Control Flow Example #4

```
0000000000400540 <multstore>:  
•  
•  
400544: call    400550 <mult2>  
400549: mov    %rax, (%rbx)   
•  
•
```

```
0000000000400550 <mult2>:  
400550:    mov        %rdi,%rax  
•  
•  
400557:    ret
```



Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - **Passing data**
 - Managing local data
- Illustration of Recursion

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

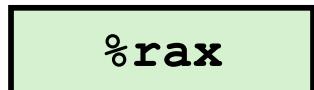
Procedure Data Flow

Registers

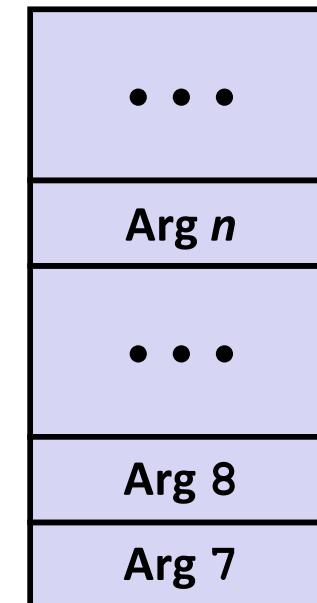
First 6 arguments



Return value



Stack



Only allocate stack space
when needed

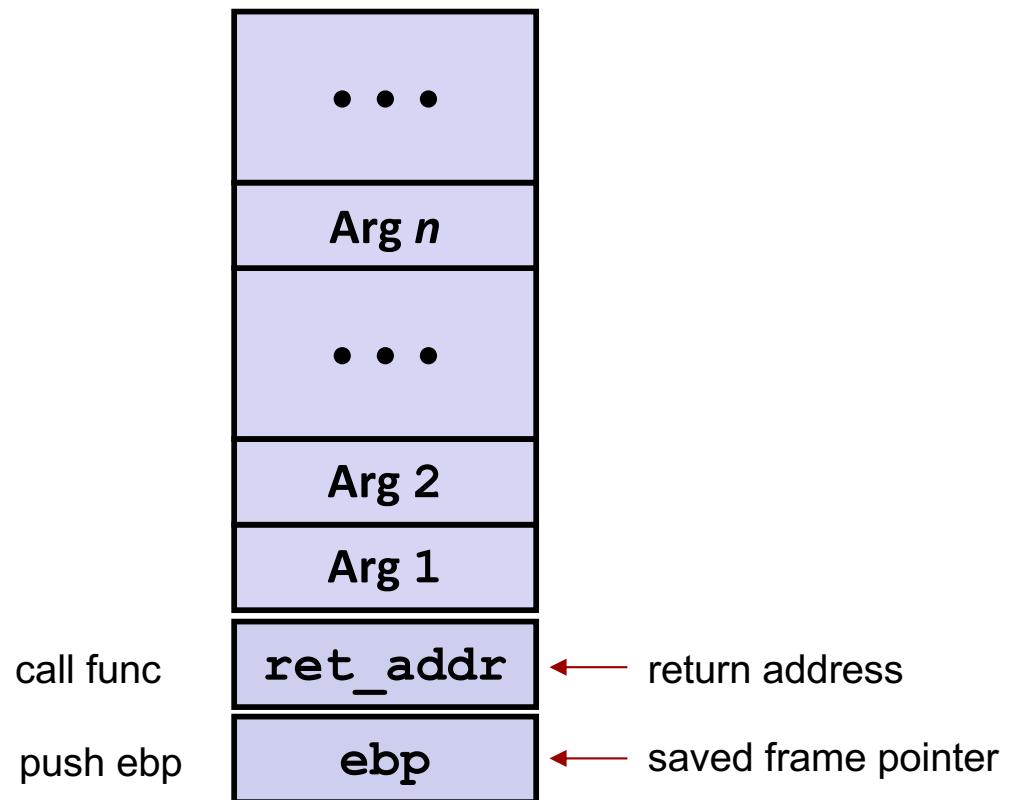
Procedure Data Flow (x86)

Registers

Return value

%eax

Stack



Procedure Data Flow (x86)

Prologue

```
push %ebp  
mov %esp,%ebp
```

Epilogue

```
mov %ebp,%esp  
pop %ebp  
ret
```

or

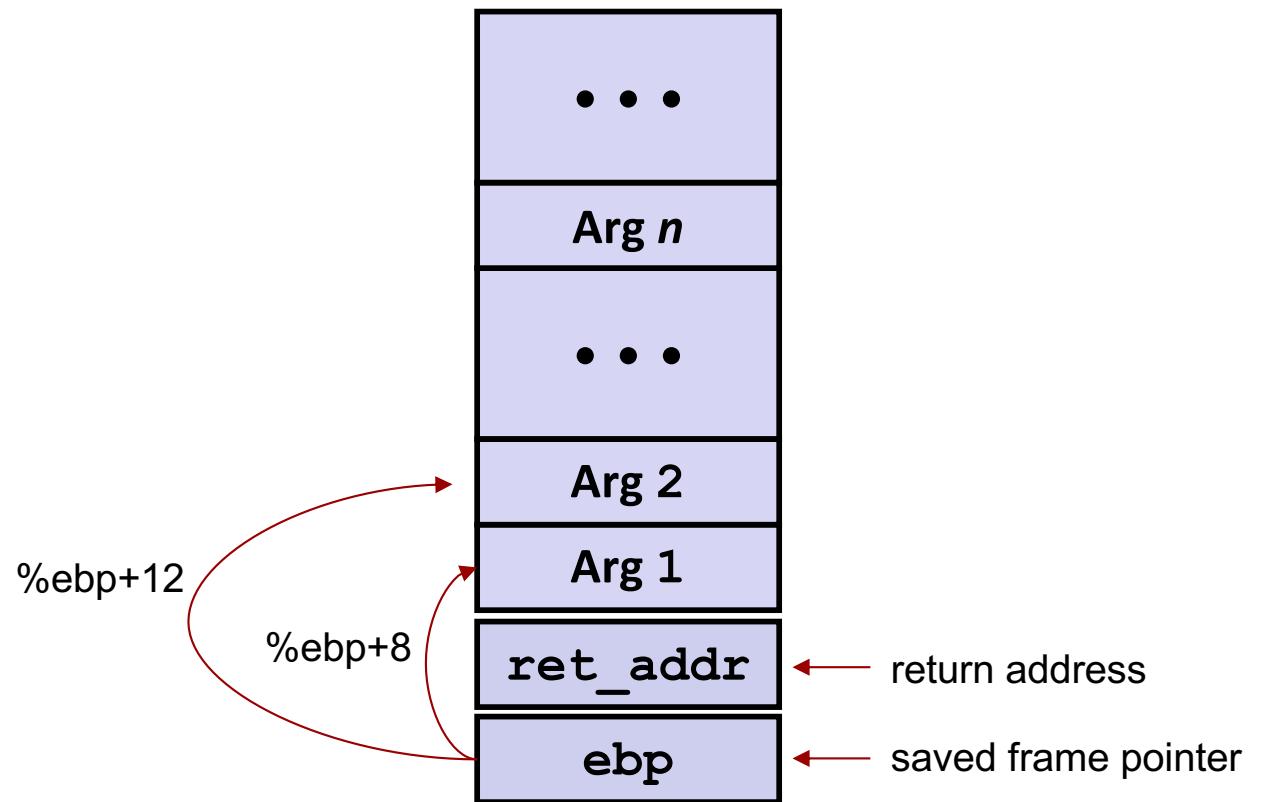
```
leave  
ret
```

Procedure Data Flow (x86)

Prologue

```
push %ebp  
mov %esp,%ebp
```

Stack



Procedure Data Flow (x86)

Example

```
#define M 9
#define N 15

int mat1[M][N];
int mat2[N][M];
int copy_element(int i, int j)
{
    mat1[i][j] = mat2[j][i];
}
```

```
copy_element:
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
    movl 8(%ebp),%ecx
    movl 12(%ebp),%ebx
    movl %ecx,%edx
    leal (%ebx,%ebx,8),%eax
    sall $4,%edx
    sall $2,%eax
    subl %ecx,%edx
    movl mat2(%eax,%ecx,4),%eax
    sall $2,%edx
    movl %eax,mat1(%edx,%ebx,4)
    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
```

Procedure Data Flow (x86)

Example

```
int loop (int x, int y)
{
    int result;
    for (result=0; x>y; result++) {
        x--;
        y++;
    }
    result++;
}
```

```
loop:
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%ecx
    movl 12(%ebp),%edx
    xorl %eax,%eax
    cmpl %edx,%ecx
    jle .L4

.L6:
    decl %ecx
    incl %edx
    incl %eax
    cmpl %edx,%ecx
    jg .L6

.L4:
    incl %eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Procedure Data Flow (x86)

Example

```
/* copy string x to buf */
void foo(char *x) {
    int buf[1];
    strcpy((char *)buf, x);
}

void callfoo() {
    foo("abcdefghijklm");
}
```

Procedure Call

Example

```
/* copy string to buf */
void foo(char *buf)
{
    int buf_size;
    strcpy(buf, "abc");
}

void callfoo()
{
    foo("abc");
}
```

080484f4 <foo>:			
080484f4: 55	pushl %ebp		
080484f5: 89 e5	movl %esp,%ebp		
080484f7: 83 ec 18	subl \$0x18,%esp		
080484fa: 8b 45 08	movl 0x8(%ebp),%eax		
080484fd: 83 c4 f8	addl \$0xffffffff8,%esp		
08048500: 50	pushl %eax		
08048501: 8d 45 fc	leal 0xfffffff(%ebp),%eax		
08048504: 50	pushl %eax		
08048505: e8 ba fe ff ff	call 80483c4 <strcpy>		
0804850a: 89 ec	movl %ebp,%esp		
0804850c: 5d	popl %ebp		
0804850d: c3	ret		
08048510 <callfoo>:			
08048510: 55	pushl %ebp		
08048511: 89 e5	movl %esp,%ebp		
08048513: 83 ec 08	subl \$0x8,%esp		
08048516: 83 c4 f4	addl \$0xffffffff4,%esp		
08048519: 68 9c 85 04 08	pushl \$0x804859c		
0804851e: e8 d1 ff ff ff	call 80484f4 <foo>		
08048523: 89 ec	movl %ebp,%esp		
08048525: 5d	popl %ebp		
08048526: c3	ret		

Procedure Data Flow (x86)

Example

```
/* copy string x to buf
void foo(char *x) {
    int buf[1];
    strcpy((char *)buf, x);
}

void callfoo() {
    foo("abcdefghijklm");
}
```

nshc@nshcdell: ~/computer_system/05_machine2

The screenshot shows the pwndbg debugger interface with the following sections:

- Registers:** Shows CPU register values. EAX is 0xfffffd4b4 (← 0x40000), EBX is 0x0, ECX is 0x3e6fc73b, EDX is 0xfffffd514 (← 0x0), EDI is 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) (← 0x1ead6c), ESI is 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) (← 0x1ead6c), EBP is 0xfffffd4b8 (→ 0xfffffd4d8 → 0xfffffd4e8 ← 0x0), ESP is 0xfffffd490 (→ 0xfffffd4b4 ← 0x40000), and EIP is 0x80491aa (foo+20) (→ 0xffffeb1e8 ← 0x0).
- Disassembly:** Shows the assembly code for the foo and callfoo functions. The foo function copies the string "abcdefghijklm" from memory at address 0x804a008 into the buffer at 0xfffffd4b4. The callfoo function calls foo.
- Stack:** Shows the current state of the stack. The stack grows downwards, with esp pointing to the top of the stack at 0xfffffd490. The stack contains the copied string "abcdefghijklm" at 0xfffffd494, the global offset table entry for _exit_funcs_lock at 0xfffffd49c, and other system library symbols.
- Backtrace:** Shows the call history of the current thread. The stack trace starts at 0x80491aa (foo+20) and goes up through 0x80491cc (callfoo+23), 0x80491e1 (main+15), and finally 0xf7de0ed5 (_libc_start_main+245).

pwndbg>

Procedure Data Flow (x86)

Example

```
/* copy string x to buf
void foo(char *x) {
    int buf[1];
    strcpy((char *)buf, x);
}

void callfoo() {
    foo("abcdefghi");
}
```

nshc@nshcdell: ~/computer_system/05_machine2

[DISASM / i386 / set emulate on]

```

0x804919d <foo+7>    sub    esp, 0x18
0x80491a0 <foo+10>   sub    esp, 8
0x80491a3 <foo+13>   push   dword ptr [ebp + 8]
0x80491a6 <foo+16>   lea     eax, [ebp - 4]
0x80491a9 <foo+19>   push   eax
► 0x80491aa <foo+20>  call    strcpy@plt      <strcpy@plt>
    dest: 0xfffffd4b4 ← 0x40000
    src: 0x804a008 ← 'abcdefghi'

0x80491af <foo+25>   add    esp, 0x10
0x80491b2 <foo+28>   nop
0x80491b3 <foo+29>   leave
0x80491b4 <foo+30>   ret

0x80491b5 <callfoo> endbr32

```

[STACK]

```

00:0000 | esp 0xfffffd490 → 0xfffffd4b4 ← 0x40000
01:0004 |          0xfffffd494 → 0x804a008 ← 'abcdefghi'
02:0008 |          0xfffffd498 → 0xf7ffc7e0 (_rtld_global_ro) ← 0x0
03:000c |          0xfffffd49c → 0xf7fb44e8 (_exit_funcs_lock) ← 0x0
04:0010 |          0xfffffd4a0 → 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) ← 0x1ead6c
05:0014 |          0xfffffd4a4 → 0xf7fe22b0 (_dl_fini) ← endbr32
06:0018 |          0xfffffd4a8 ← 0x0
07:001c |          0xfffffd4ac → 0xf7dfa352 (_internal_atexit+66) ← add esp, 0x10

```

[BACKTRACE]

```

► 0 0x80491aa foo+20
1 0x80491cc callfoo+23
2 0x80491e1 main+15
3 0xf7de0ed5 __libc_start_main+245

```

pwndbg> x/32xb 0xfffffd4b4

0xfffffd4b4:	0x00	0x00	0x04	0x00	0xd8	0xd4	0xff	0xff
0xfffffd4bc:	0xcc	0x91	0x04	0x08	0x08	0xa0	0x04	0x08
0xfffffd4c4:	0x84	0xa5	0x7f	0x7f	0x8c	0xd5	0xff	0xff
0xfffffd4cc:	0x11	0x92	0x04	0x08	0xb0	0x22	0xfe	0xf7

pwndbg>

Procedure Data Flow (x86)

Example

```
/* copy string x to b
void foo(char *x) {
    int buf[1];
    strcpy((char *)bu
}

void callfoo() {
    foo("abcdefghi");
}
```

nshc@nshcdell: ~/computer_system/05_machine2

*ECX 0x804a008 ← 'abcdefghi'
*EDX 0xfffffd4b4 ← 'abcdefghi'
EDI 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) ← 0x1ead6c
ESI 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) ← 0x1ead6c
EBP 0xfffffd4b8 ← 'efghi'
ESP 0xfffffd490 → 0xfffffd4b4 ← 'abcdefghi'
*EIP 0x80491af (foo+25) ← add esp, 0x10

[DISASM / i386 / set emulate on]

0x80491a0 <foo+10>	sub	esp, 8
0x80491a3 <foo+13>	push	dword ptr [ebp + 8]
0x80491a6 <foo+16>	lea	eax, [ebp - 4]
0x80491a9 <foo+19>	push	eax
0x80491aa <foo+20>	call	strcpy@plt

<strcpy@plt>

► 0x80491af <foo+25> add esp, 0x10

0x80491b2 <foo+28>	nop
0x80491b3 <foo+29>	leave
0x80491b4 <foo+30>	ret

0x80491b5 <callfoo> endbr32
0x80491b9 <callfoo+4> push ebp

[STACK]

00:0000 esp 0xfffffd490 → 0xfffffd4b4 ← 'abcdefghi'
01:0004 0xfffffd494 → 0x804a008 ← 'abcdefghi'
02:0008 0xfffffd498 → 0xf7ffc7e0 (_rtld_global_ro) ← 0x0
03:000c 0xfffffd49c → 0xf7fb44e8 (_exit_funcs_lock) ← 0x0
04:0010 0xfffffd4a0 → 0xf7fb1000 (_GLOBAL_OFFSET_TABLE_) ← 0x1ead6c
05:0014 0xfffffd4a4 → 0xf7fe22b0 (_dl_fini) ← endbr32
06:0018 0xfffffd4a8 ← 0x0
07:001c 0xfffffd4ac → 0xf7dfa352 (_internal_atexit+66) ← add esp, 0x10

[BACKTRACE]

► 0 0x80491af foo+25
1 0x8040069
2 0x804a008

pwndbg> x/32xb 0xfffffd4b4

0xfffffd4b4:	0x61	0x62	0x62	0x64	0x65	0x66	0x67	0x68
0xfffffd4bc:	0x69	0x00	0x04	0x08	0x08	0xa0	0x04	0x08
0xfffffd4c4:	0x84	0xd5	0x1f	0x1f	0x8c	0xd5	0xff	0xff
0xfffffd4cc:	0x11	0x92	0x04	0x08	0xb0	0x22	0xfe	0xf7

pwndbg>

Data Flow Examples

```
void multstore
    (long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
0000000000400540 <multstore>:
# x in %rdi, y in %rsi, dest in %rdx
• • •
400541: mov    %rdx,%rbx          # Save dest
400544: call   400550 <mult2>      # mult2(x,y)
# t in %rax
400549: mov    %rax,(%rbx)        # Save at dest
• • •
```

```
long mult2
    (long a, long b)
{
    long s = a * b;
    return s;
}
```

```
0000000000400550 <mult2>:
# a in %rdi, b in %rsi
400550: mov    %rdi,%rax          # a
400553: imul   %rsi,%rax          # a * b
# s in %rax
400557: ret                   # Return
```

Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- Illustration of Recursion

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

Stack-Based Languages

Languages that support recursion

- e.g., C, Pascal, Java
- Code must be “*Reentrant*”
 - Multiple simultaneous instantiations of single procedure
- Need some place to store state of each instantiation
 - Arguments
 - Local variables
 - Return pointer

Stack discipline

- State for given procedure needed for limited time
 - From when called to when return
- Callee returns before caller does

Stack allocated in *Frames*

- state for single procedure instantiation

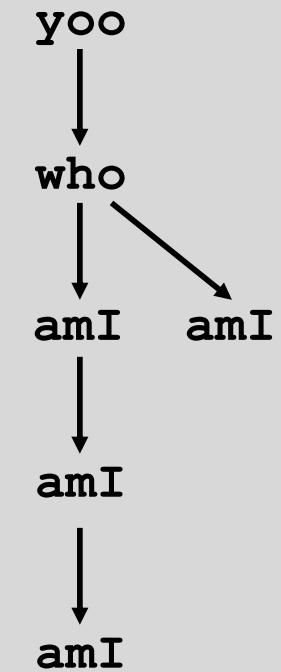
Call Chain Example

```
yoo (...)  
{  
    •  
    •  
    who () ;  
    •  
    •  
}
```

```
who (...)  
{  
    • • •  
    amI () ;  
    • • •  
    amI () ;  
    • • •  
}
```

```
amI (...)  
{  
    •  
    •  
    amI () ;  
    •  
    •  
}
```

Example Call Chain

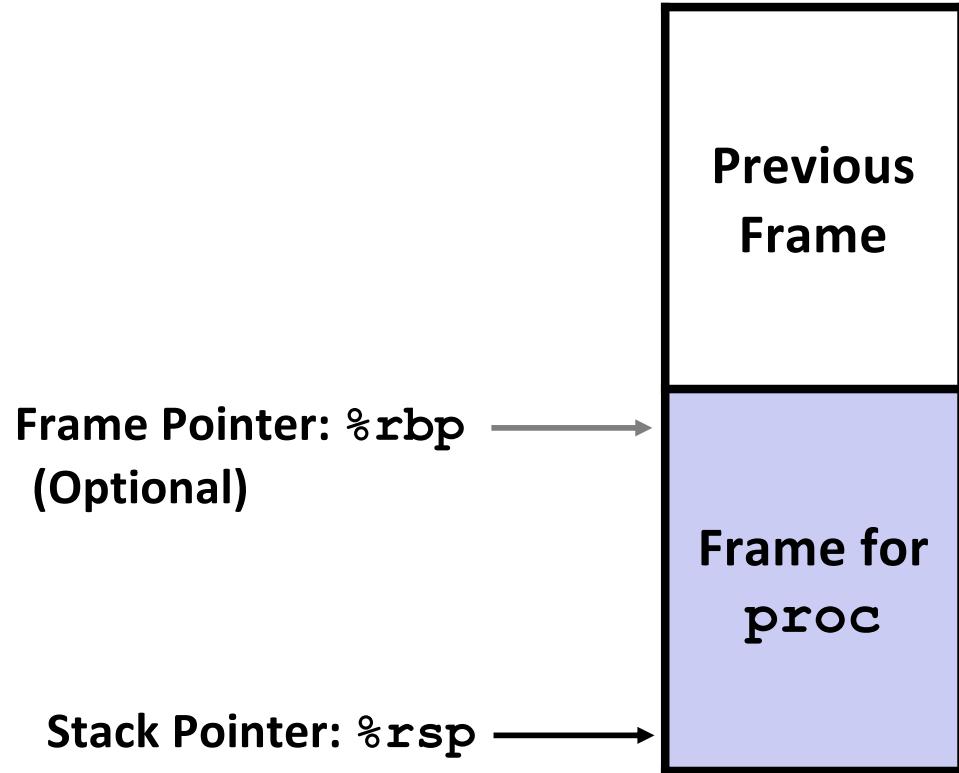


Procedure **amI ()** is recursive

Stack Frames

Contents

- Return information
- Local storage (if needed)
- Temporary space (if needed)

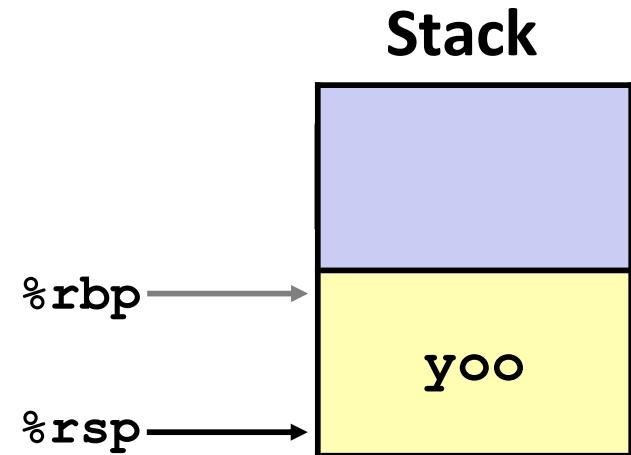
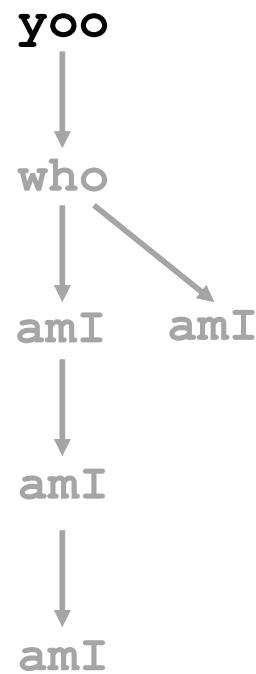


Management

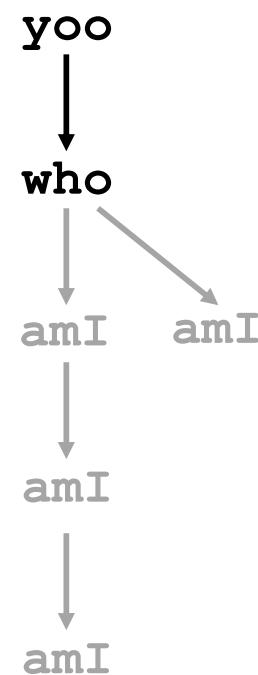
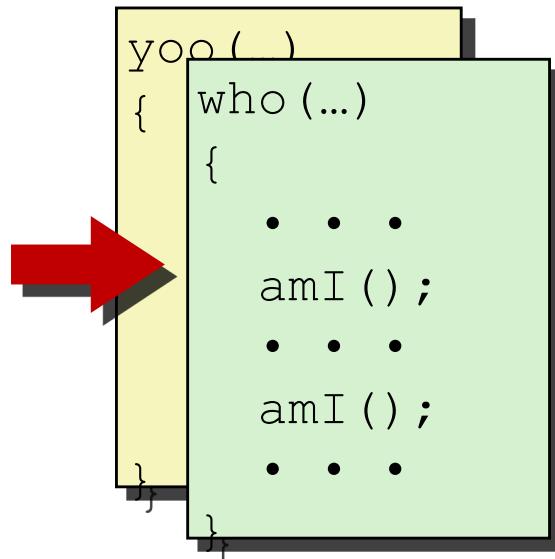
- Space allocated when enter procedure
 - “Set-up” code
 - Includes push by **call** instruction
- Deallocated when return
 - “Finish” code
 - Includes pop by **ret** instruction

Example

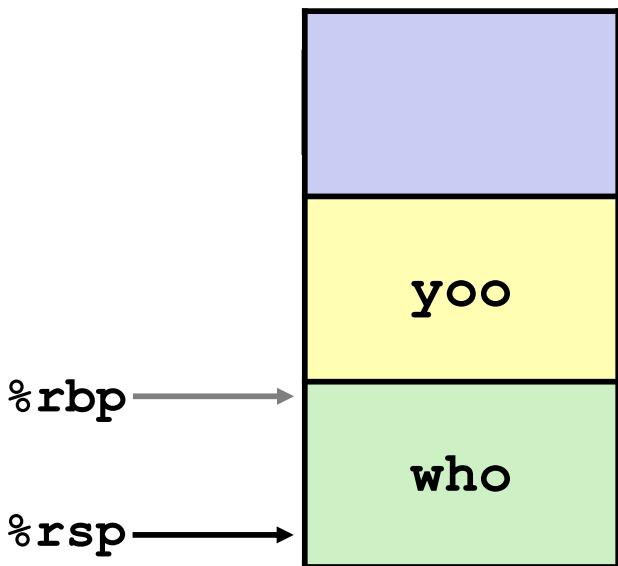
```
yoo (...)  
{  
    •  
    •  
    who () ;  
    •  
    •  
}
```



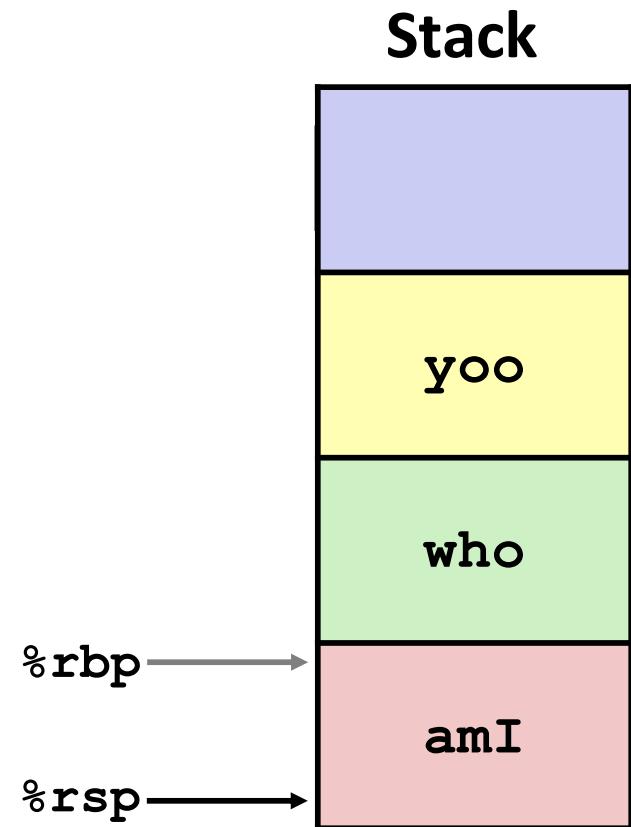
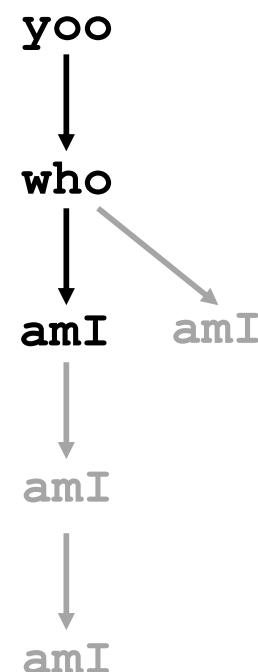
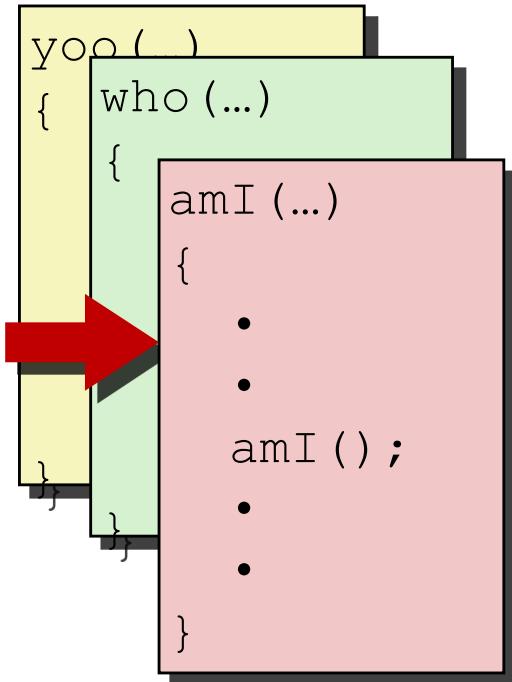
Example



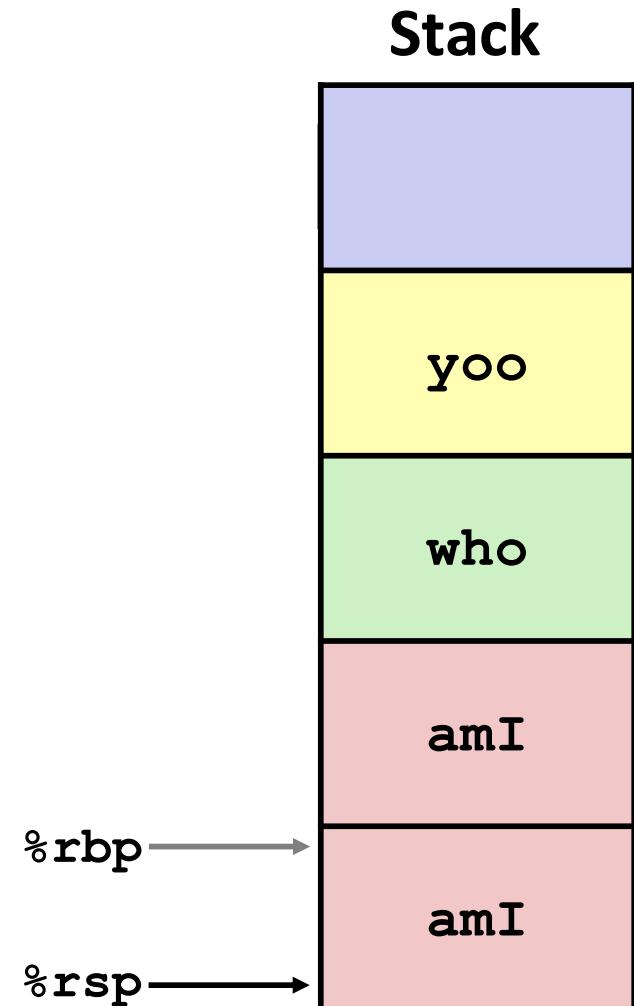
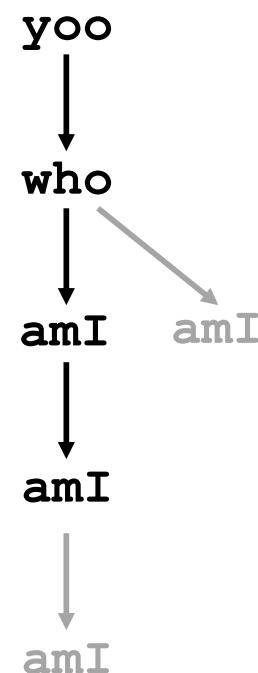
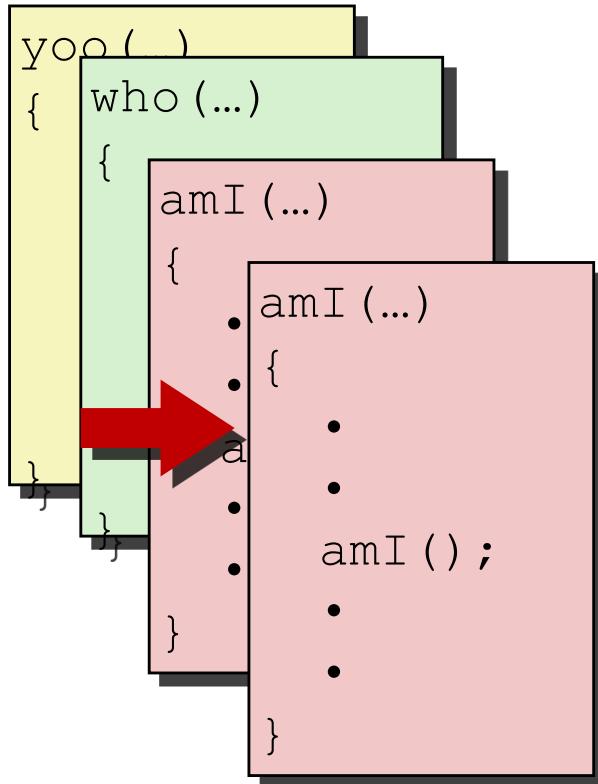
Stack



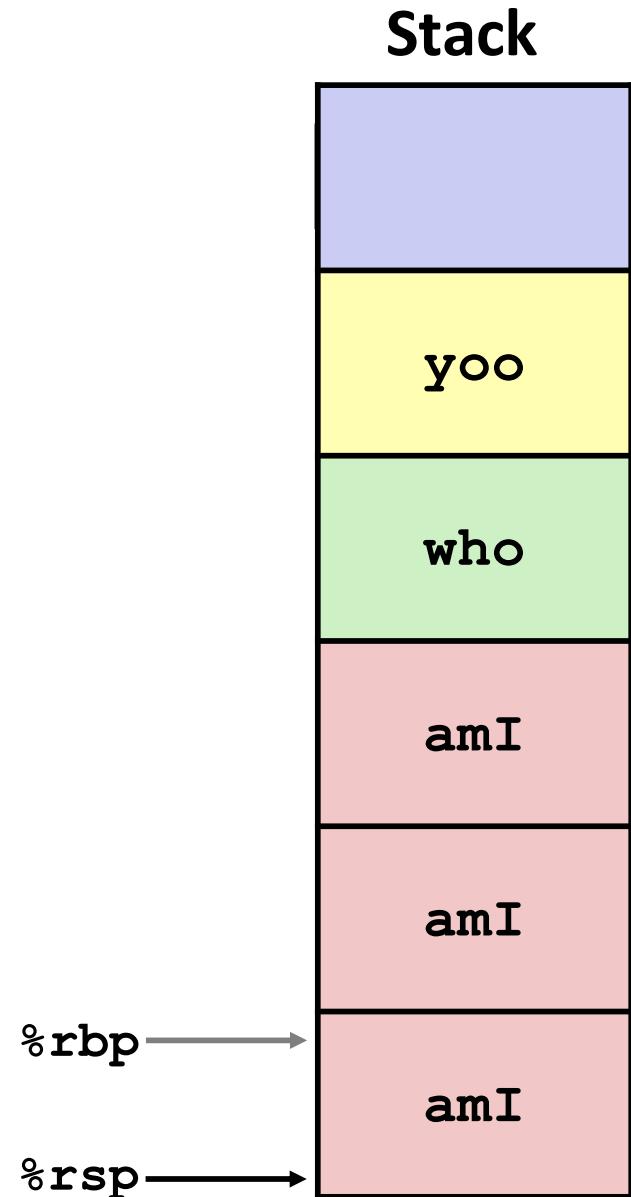
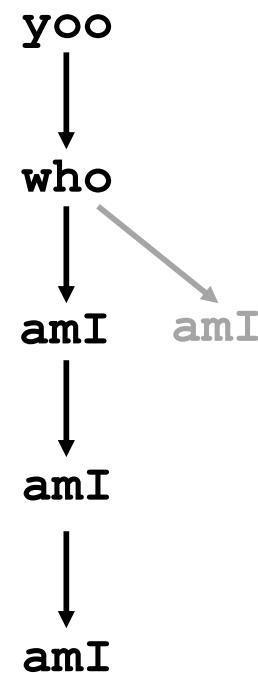
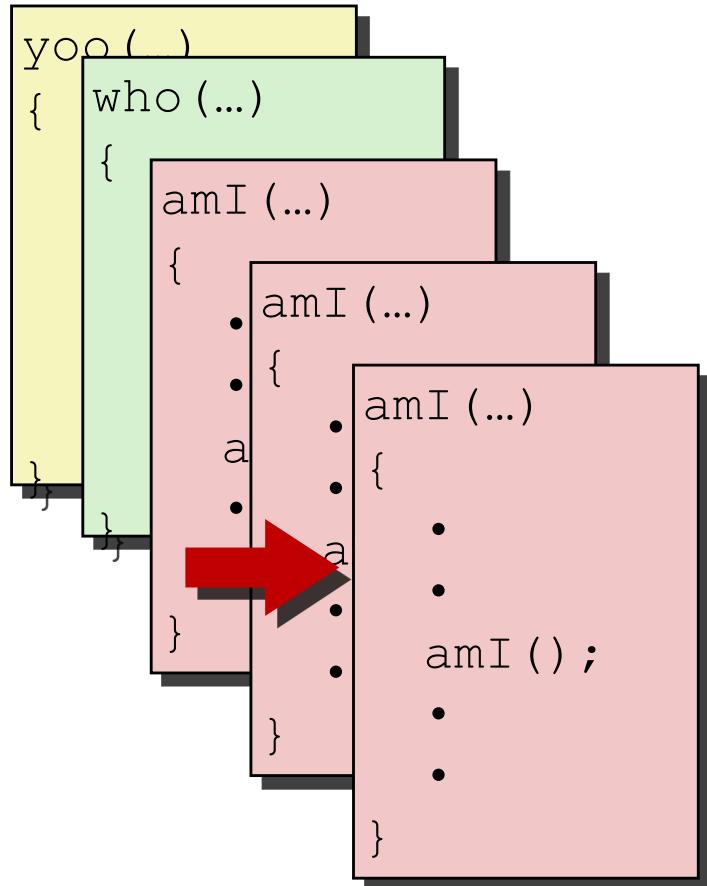
Example



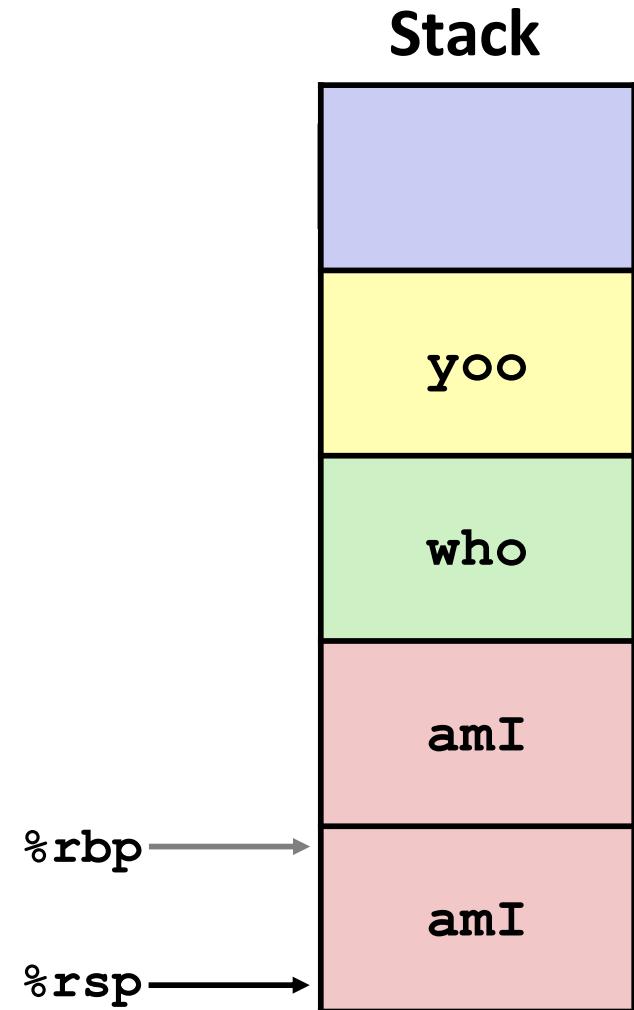
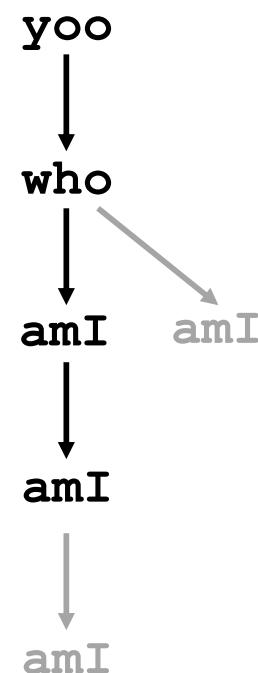
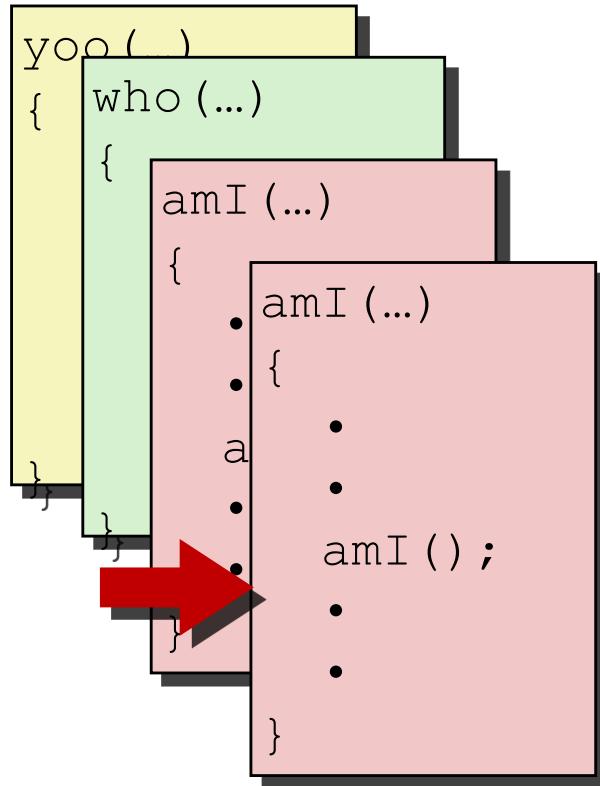
Example



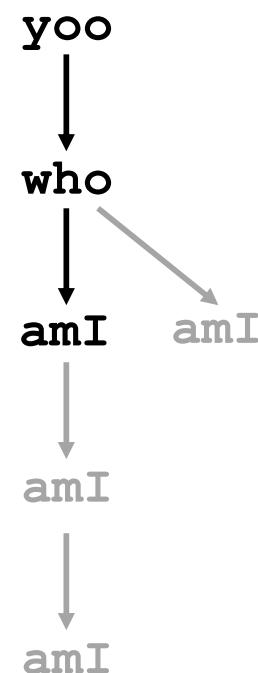
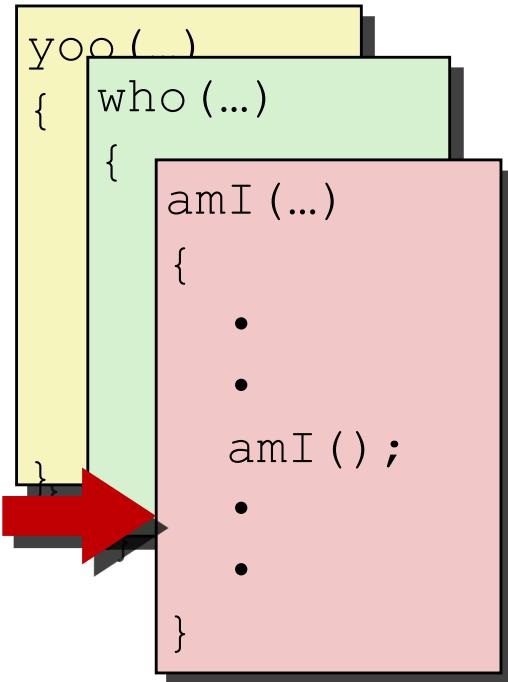
Example



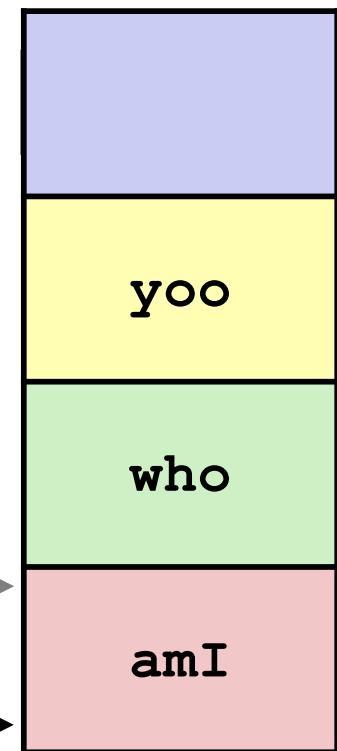
Example



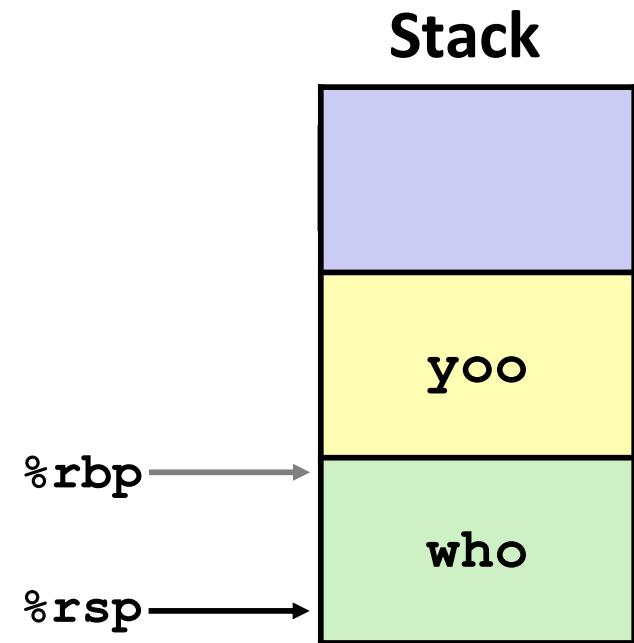
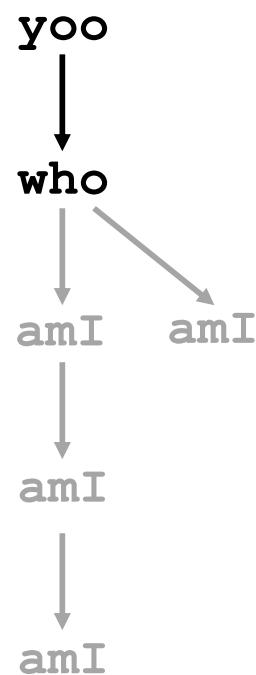
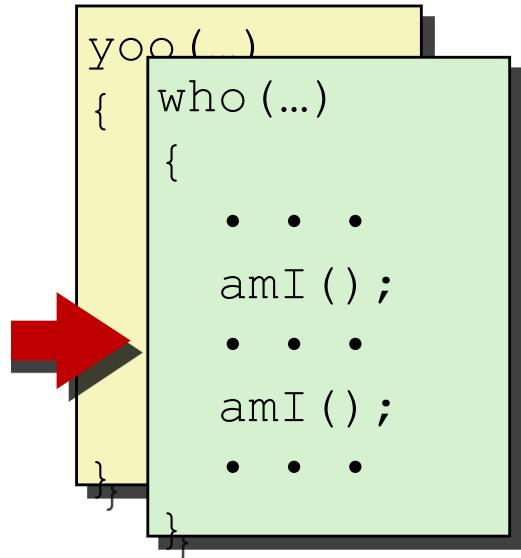
Example



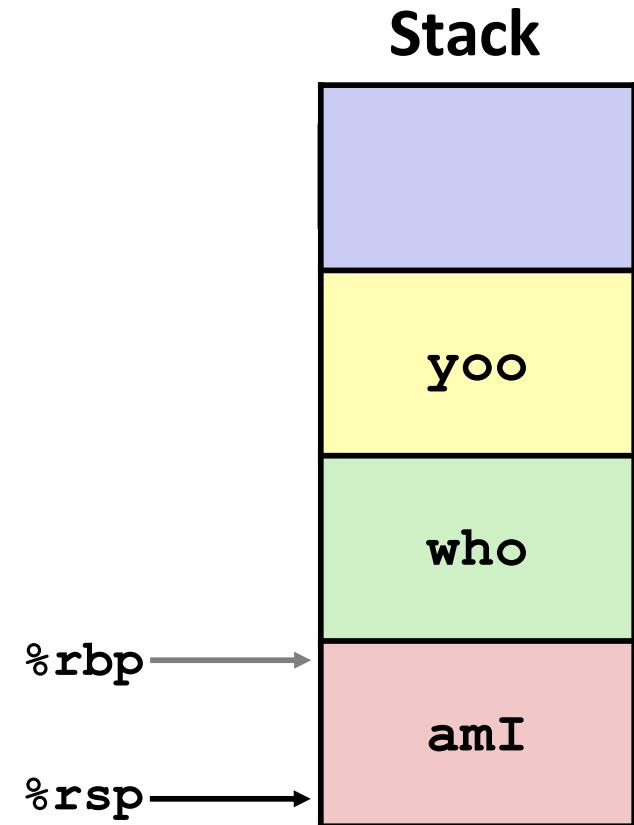
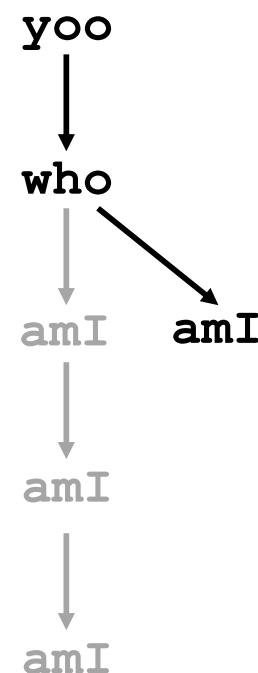
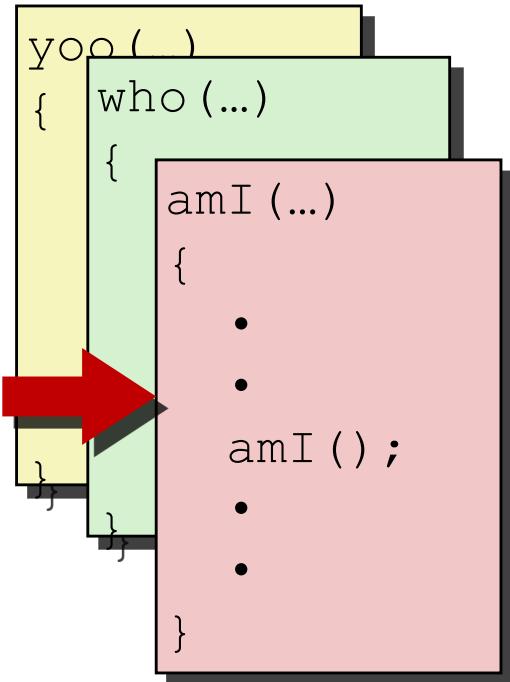
Stack



Example

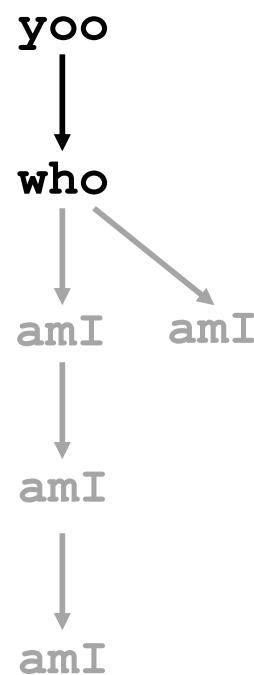


Example

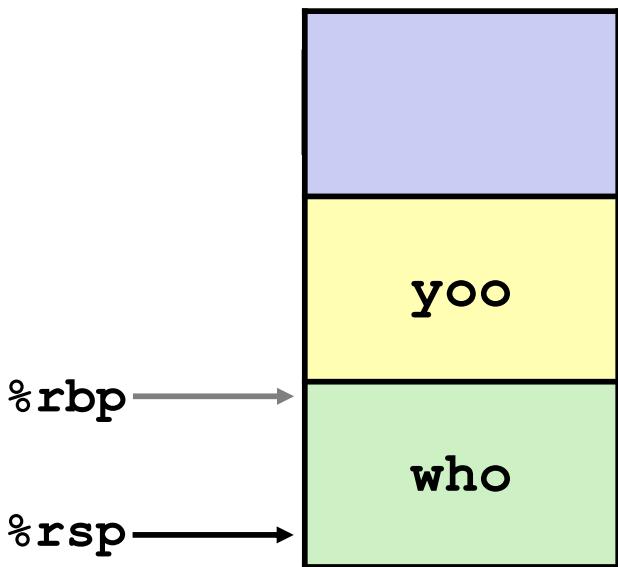


Example

```
yoo( )  
{   who (...)  
{  
    . . .  
    amI () ;  
    . . .  
    amI () ;  
    . . .  
}
```

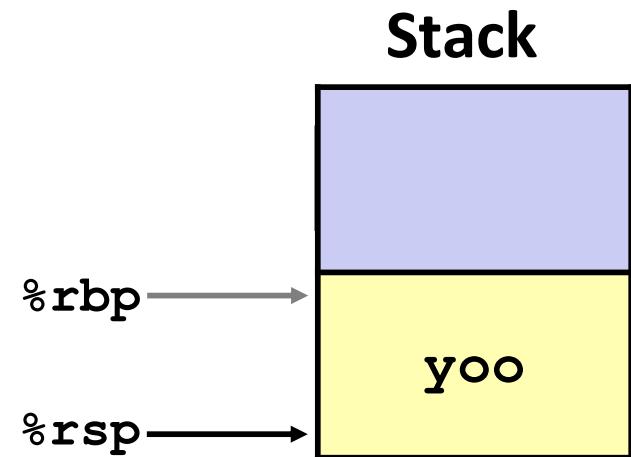
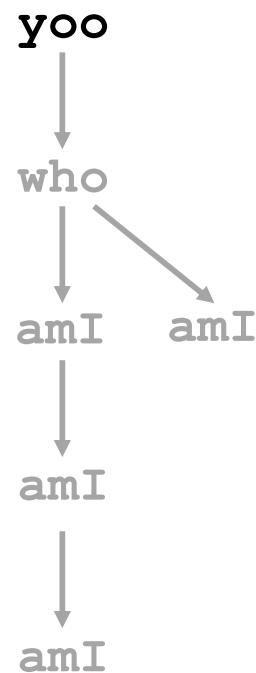


Stack



Example

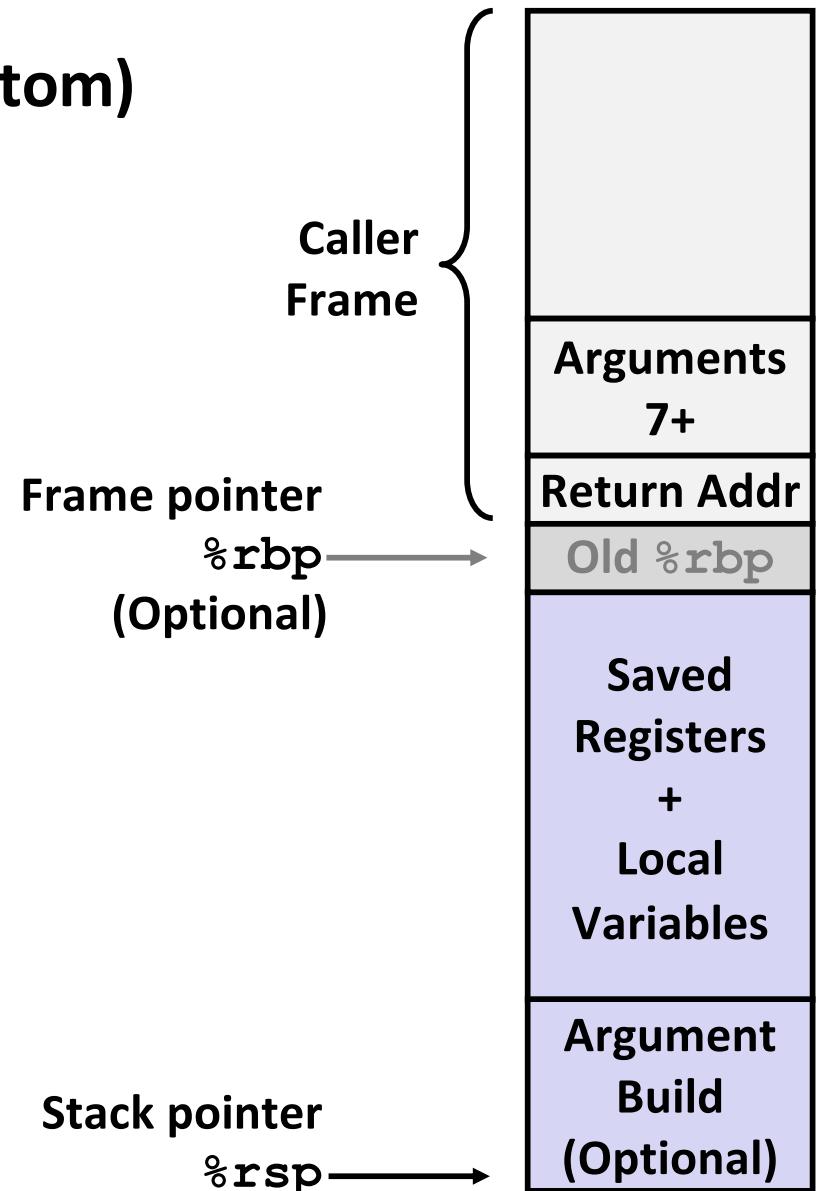
```
yoo (...)  
{  
    •  
    •  
    who () ;  
    •  
    •  
}  
}
```



x86-64/Linux Stack Frame

Current Stack Frame (“Top” to Bottom)

- “Argument build:”
Parameters for function about to call
- Local variables
If can’t keep in registers
- Saved register context
- Old frame pointer (optional)



Caller Stack Frame

- Return address
 - Pushed by **call** instruction
- Arguments for this call

Example: incr

```
long incr(long *p, long val) {  
    long x = *p;  
    long y = x + val;  
    *p = y;  
    return x;  
}
```

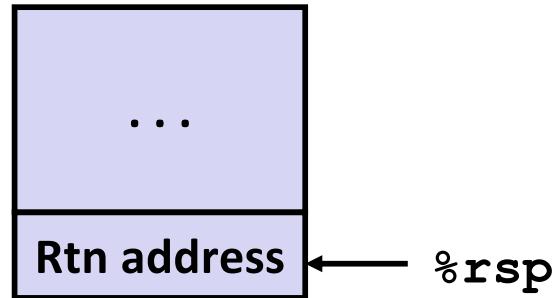
```
incr:  
    movq    (%rdi), %rax  
    addq    %rax, %rsi  
    movq    %rsi, (%rdi)  
    ret
```

Register	Use(s)
%rdi	Argument p
%rsi	Argument val , y
%rax	x , Return value

Example: Calling `incr #1`

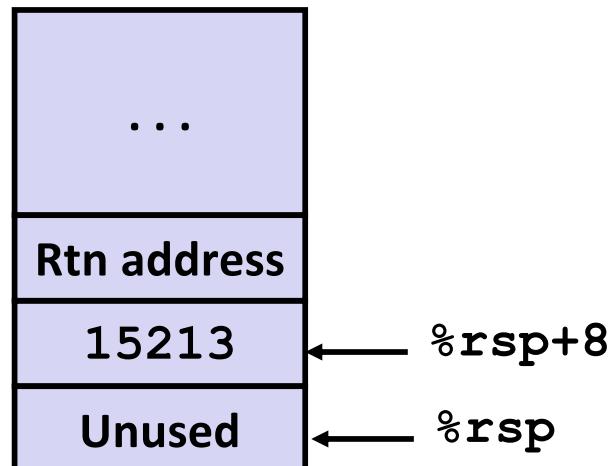
```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

Initial Stack Structure



```
call_incr:
    subq    $16, %rsp
    movq    $15213, 8(%rsp)
    movl    $3000, %esi
    leaq    8(%rsp), %rdi
    call    incr
    addq    8(%rsp), %rax
    addq    $16, %rsp
    ret
```

Resulting Stack Structure

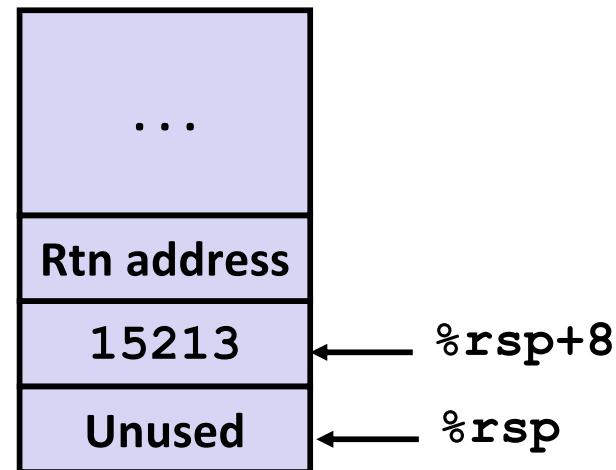


Example: Calling `incr` #2

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

```
call_incr:
    subq    $16, %rsp
    movq    $15213, 8(%rsp)
    movl    $3000, %esi
    leaq    8(%rsp), %rdi
    call    incr
    addq    8(%rsp), %rax
    addq    $16, %rsp
    ret
```

Stack Structure

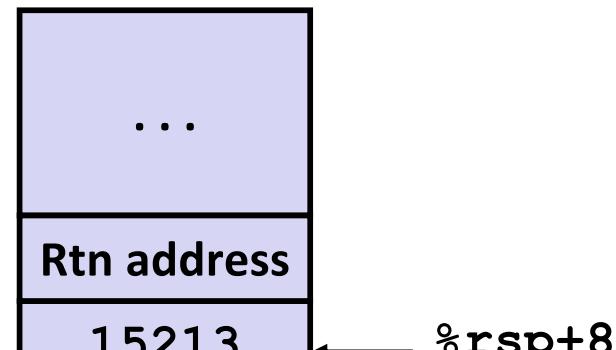


Register	Use(s)
<code>%rdi</code>	<code>&v1</code>
<code>%rsi</code>	<code>3000</code>

Example: Calling `incr` #2

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

Stack Structure



Aside 1: `movl $3000, %esi`

- ca
- Remember, `movl` -> `%exx` zeros out high order 32 bits.
 - Why use `movl` instead of `movq`? 1 byte shorter.

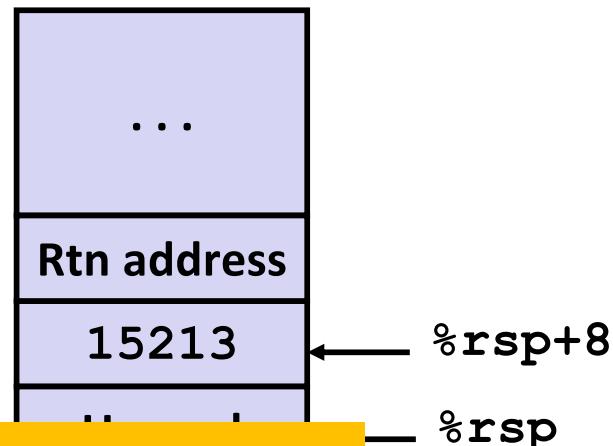
```
movl    $3000, %esi
leaq    8(%rsp), %rdi
call    incr
addq    8(%rsp), %rax
addq    $16, %rsp
ret
```

<code>%rdi</code>	<code>&v1</code>
<code>%rsi</code>	3000

Example: Calling `incr #2`

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

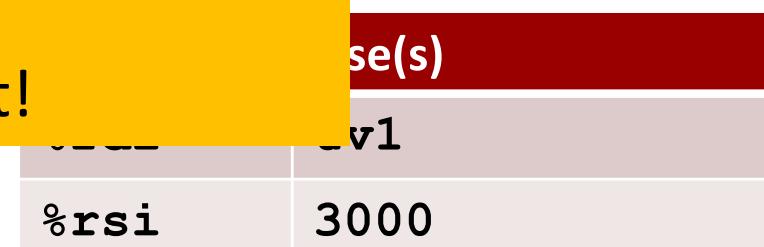
Stack Structure



call_incr
Aside 2: `leaq 8(%rsp), %rdi`

- Computes %rsp+8
- Actually, used for what it is meant!

```
leaq    8(%rsp), %rdi
call    incr
addq    8(%rsp), %rax
addq    $16, %rsp
ret
```

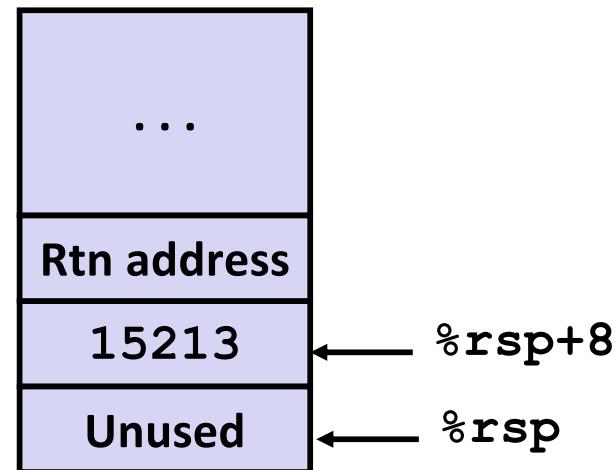


Example: Calling `incr` #2

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

```
call_incr:
    subq    $16, %rsp
    movq    $15213, 8(%rsp)
    movl    $3000, %esi
    leaq    8(%rsp), %rdi
    call    incr
    addq    8(%rsp), %rax
    addq    $16, %rsp
    ret
```

Stack Structure



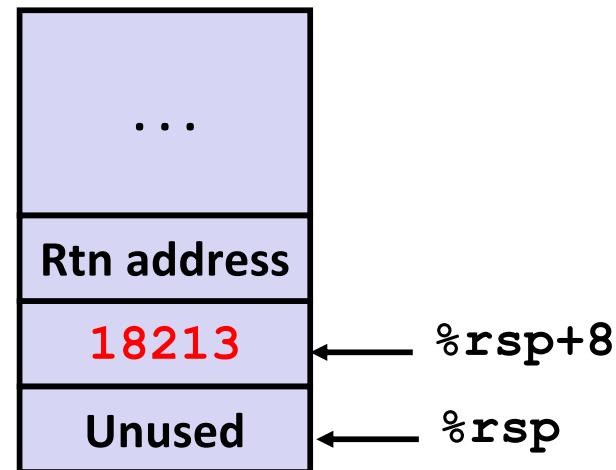
Register	Use(s)
%rdi	&v1
%rsi	3000

Example: Calling `incr` #3

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

```
call_incr:
    subq    $16, %rsp
    movq    $15213, 8(%rsp)
    movl    $3000, %esi
    leaq    8(%rsp), %rdi
    call    incr
    addq    8(%rsp), %rax
    addq    $16, %rsp
    ret
```

Stack Structure

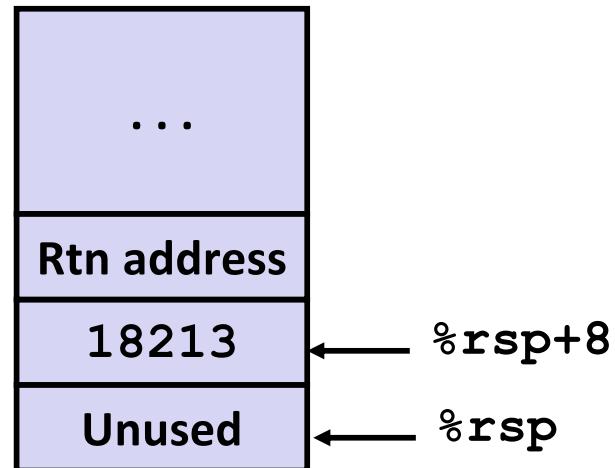


Register	Use(s)
%rdi	&v1
%rsi	3000

Example: Calling `incr` #4

Stack Structure

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```



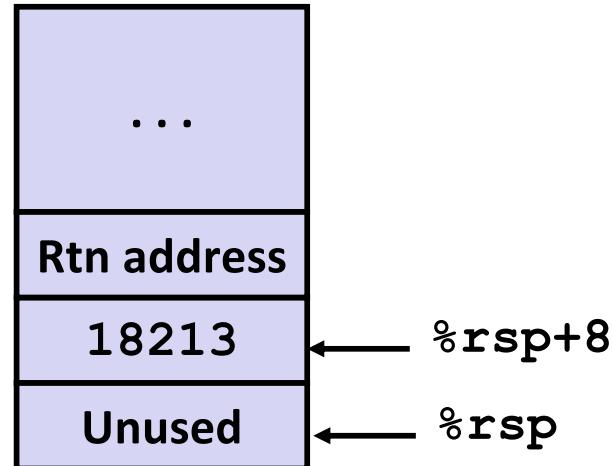
```
call_incr:
    subq    $16, %rsp
    movq    $15213, 8(%rsp)
    movl    $3000, %esi
    leaq    8(%rsp), %rdi
    call    incr
    addq    8(%rsp), %rax
    addq    $16, %rsp
    ret
```

Register	Use(s)
%rax	Return value

Example: Calling `incr` #5a

Stack Structure

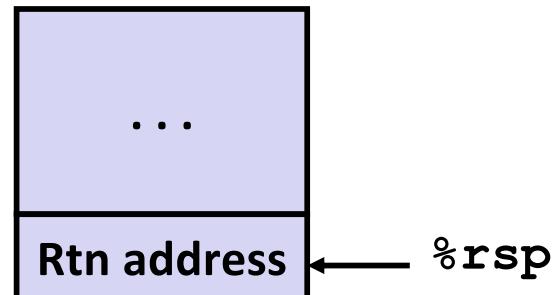
```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```



```
call_incr:
subq    $16, %rsp
movq    $15213, 8(%rsp)
movl    $3000, %esi
leaq    8(%rsp), %rdi
call    incr
addq    8(%rsp), %rax
addq    $16, %rsp
ret
```

Register	Use(s)
%rax	Return value

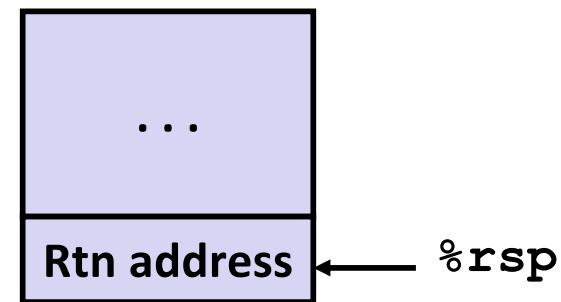
Updated Stack Structure



Example: Calling `incr #5b`

```
long call_incr() {
    long v1 = 15213;
    long v2 = incr(&v1, 3000);
    return v1+v2;
}
```

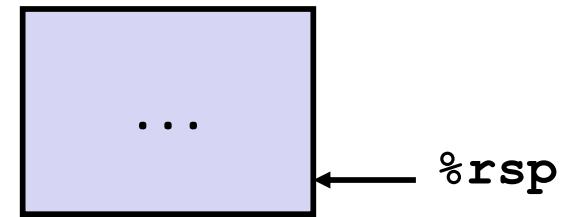
Updated Stack Structure



```
call_incr:
subq    $16, %rsp
movq    $15213, 8(%rsp)
movl    $3000, %esi
leaq    8(%rsp), %rdi
call    incr
addq    8(%rsp), %rax
addq    $16, %rsp
ret
```

Register	Use(s)
%rax	Return value

Final Stack Structure



Register Saving Conventions

When procedure **yoo** calls **who**:

- **yoo** is the *caller*
- **who** is the *callee*

Can register be used for temporary storage?

```
yoo:
```

```
    • • •  
    movq $15213, %rdx  
    call who  
    addq %rdx, %rax  
    • • •  
    ret
```

```
who:
```

```
    • • •  
    subq $18213, %rdx  
    • • •  
    ret
```

- Contents of register **%rdx** overwritten by **who**
- This could be trouble → something should be done!
 - Need some coordination

Register Saving Conventions

When procedure **yoo** calls **who**:

- **yoo** is the *caller*
- **who** is the *callee*

Can register be used for temporary storage?

Conventions

- ***“Caller Saved” (aka “Call-Clobbered”)***
 - Caller saves temporary values in its frame before the call
- ***“Callee Saved” (aka “Call-Preserved”)***
 - Callee saves temporary values in its frame before using
 - Callee restores them before returning to caller

x86-64 Linux Register Usage #1

%rax

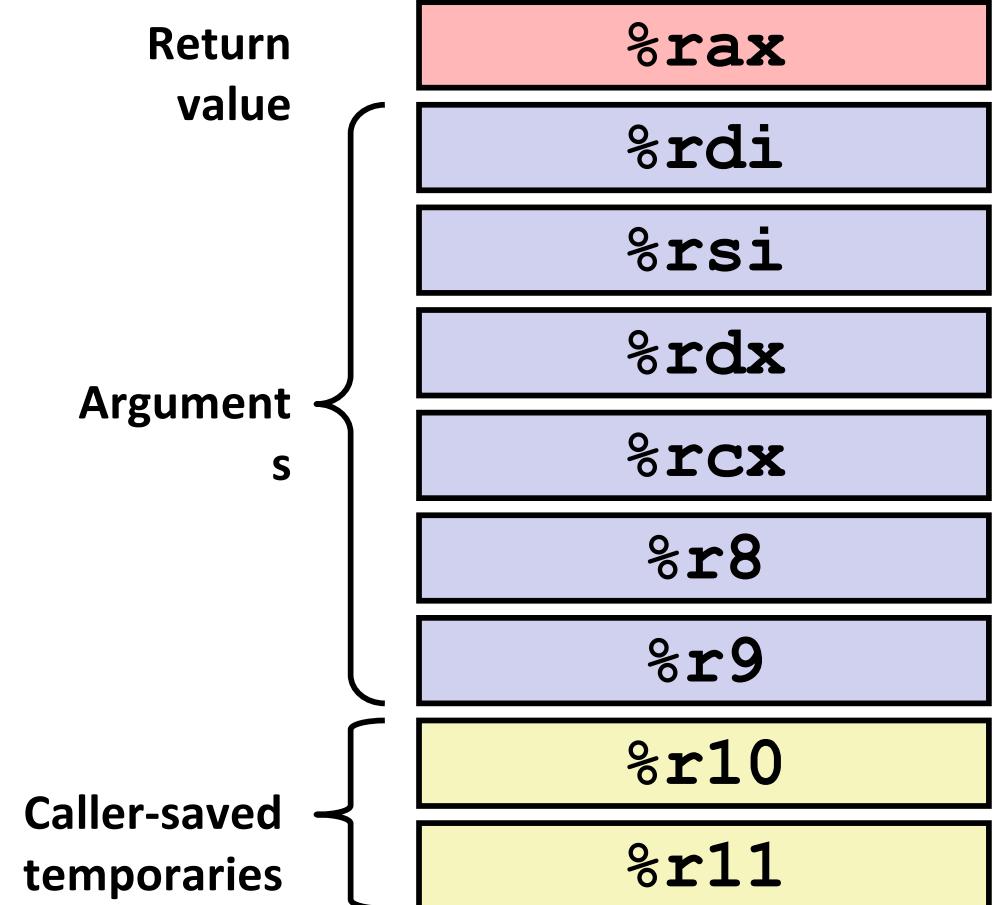
- Return value
- Also caller-saved
- Can be modified by procedure

%rdi, ..., %r9

- Arguments
- Also caller-saved
- Can be modified by procedure

%r10, %r11

- Caller-saved
- Can be modified by procedure



x86-64 Linux Register Usage #2

%rbx, %r12, %r13, %r14

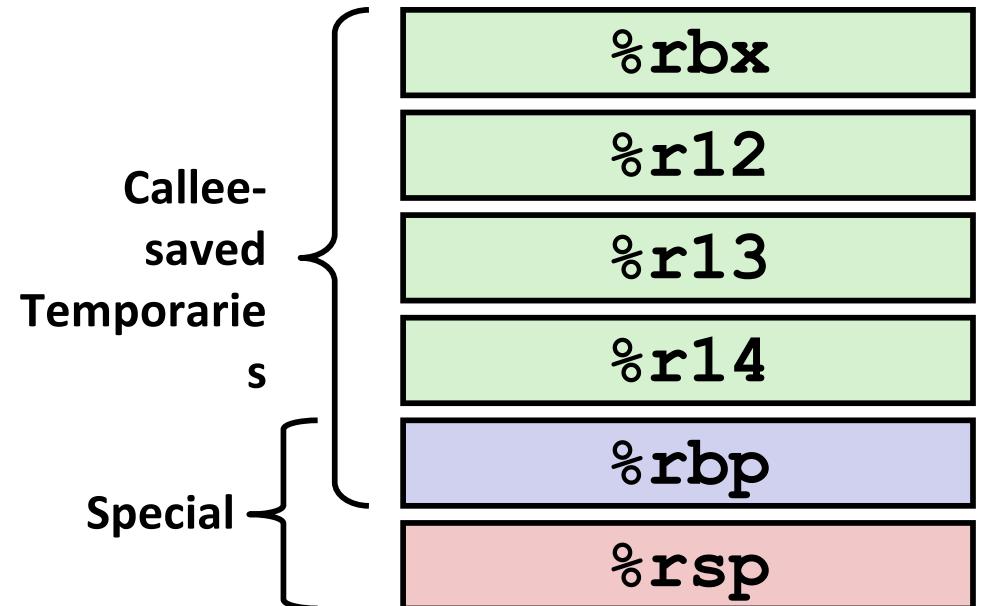
- Callee-saved
- Callee must save & restore

%rbp

- Callee-saved
- Callee must save & restore
- May be used as frame pointer
- Can mix & match

%rsp

- Special form of callee save
- Restored to original value upon exit from procedure



Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- Illustration of Recursion

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

Recursive Function

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

pcount_r:

```
    movl    $0, %eax
    testq   %rdi, %rdi
    je      .L6
    pushq   %rbx
    movq   %rdi, %rbx
    andl   $1, %ebx
    shrq   %rdi
    call   pcount_r
    addq   %rbx, %rax
    popq   %rbx
```

.L6:

```
    rep; ret
```

Recursive Function Terminal Case

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rdi	x	Argument
%rax	Return value	Return value

Recursive Function Register Save

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

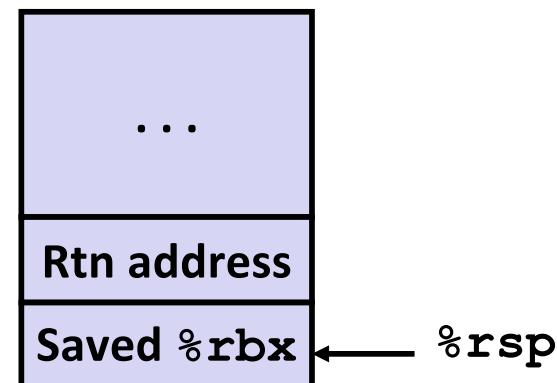
pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rdi	x	Argument



Recursive Function Call Setup

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rdi	x >> 1	Rec. argument
%rbx	x & 1	Callee-saved

Recursive Function Call

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rbx	x & 1	Callee-saved
%rax	Recursive call return value	

Recursive Function Result

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rbx	x & 1	Callee-saved
%rax	Return value	

Recursive Function Completion

```
/* Recursive popcount */
long pcount_r(unsigned long x) {
    if (x == 0)
        return 0;
    else
        return (x & 1)
            + pcount_r(x >> 1);
}
```

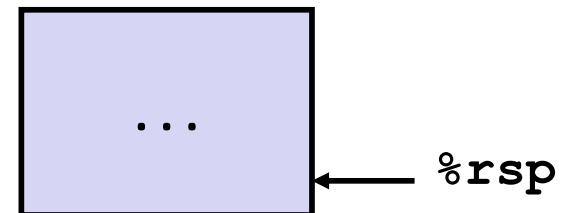
pcount_r:

movl	\$0, %eax
testq	%rdi, %rdi
je	.L6
pushq	%rbx
movq	%rdi, %rbx
andl	\$1, %ebx
shrq	%rdi
call	pcount_r
addq	%rbx, %rax
popq	%rbx

.L6:

rep; ret

Register	Use(s)	Type
%rax	Return value	Return value



Observations About Recursion

Handled Without Special Consideration

- Stack frames mean that each function call has private storage
 - Saved registers & local variables
 - Saved return pointer
- Register saving conventions prevent one function call from corrupting another's data
 - Unless the C code explicitly does so (e.g., buffer overflow in Lecture 9)
- Stack discipline follows call / return pattern
 - If P calls Q, then Q returns before P
 - Last-In, First-Out

Also works for mutual recursion

- P calls Q; Q calls P

x86-64 Procedure Summary

Important Points

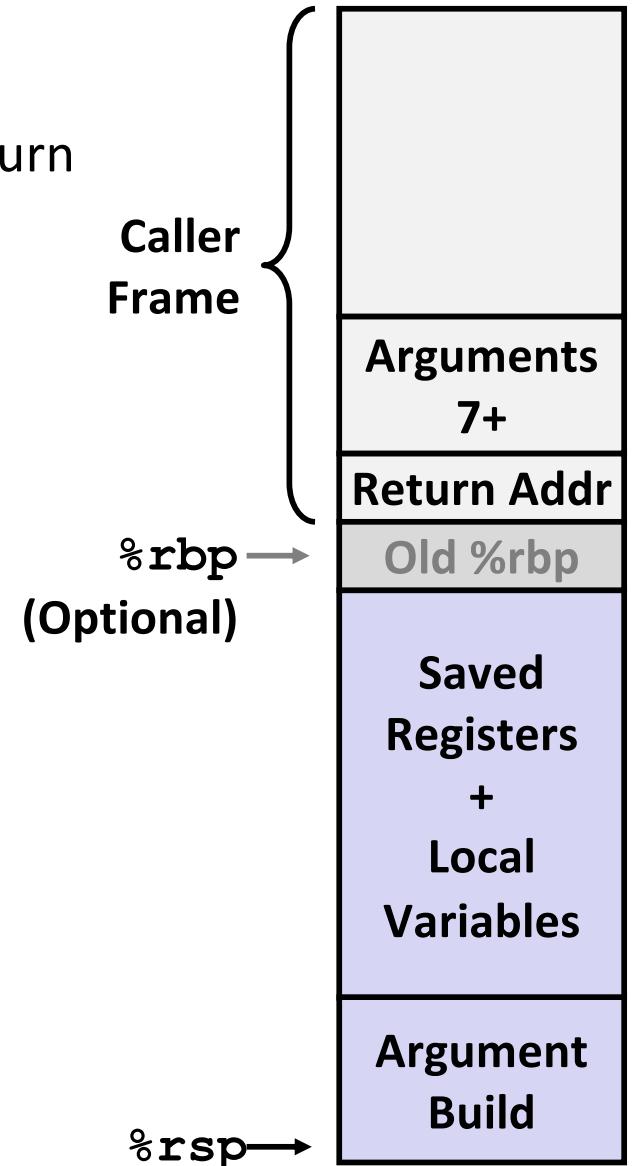
- Stack is the right data structure for procedure call/return
 - If P calls Q, then Q returns before P

Recursion (& mutual recursion) handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result return in `%rax`

Pointers are addresses of values

- On stack or global



Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- Illustration of Recursion

Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

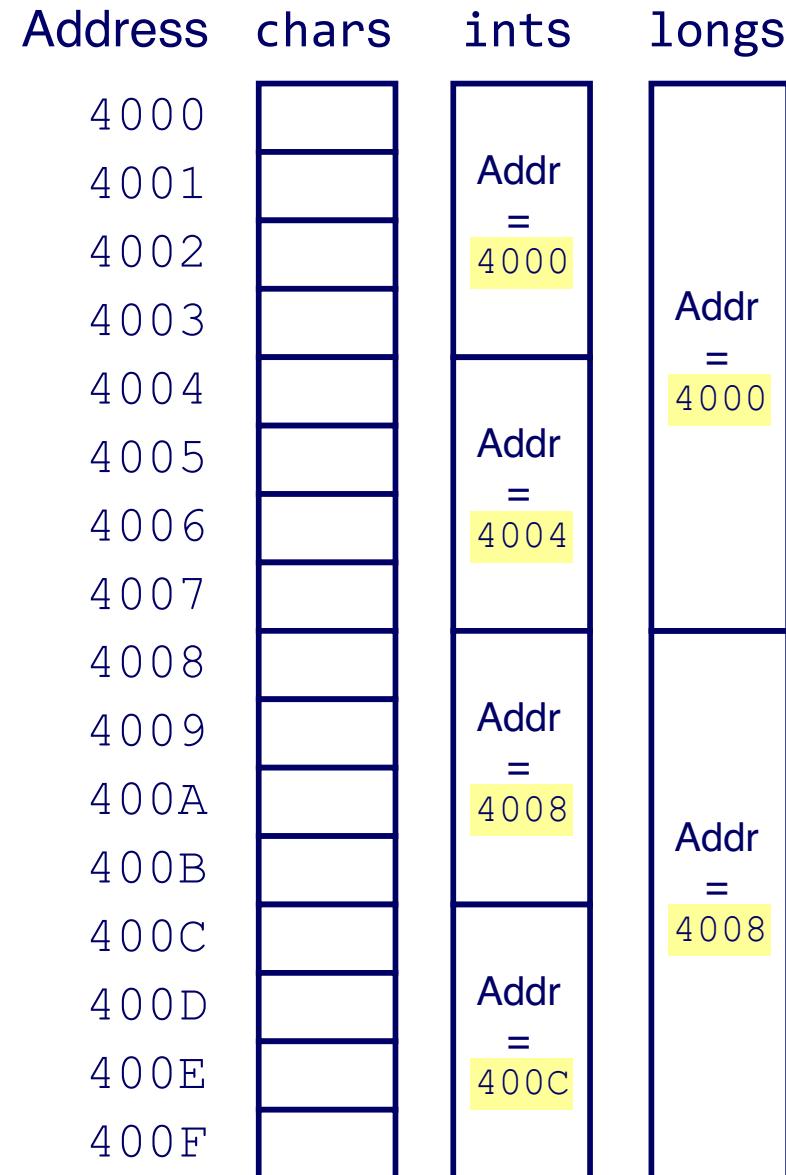
Reminder: Memory Organization

Memory locations do not have data types

- Types are implicit in how machine instructions *use* memory

Addresses specify byte locations

- Address of a larger datum is the address of its first byte
- Addresses of successive items differ by the item's size

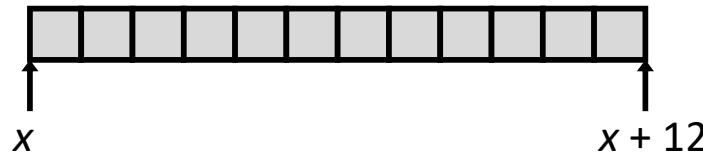


Array Allocation

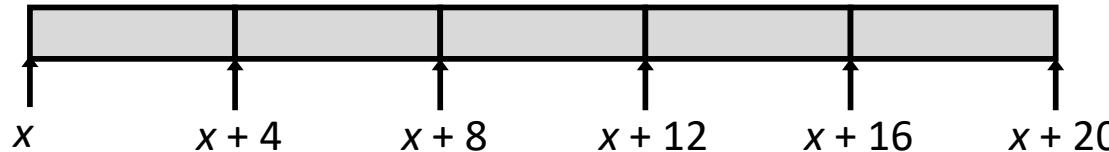
C declaration **Type name[Length]** ;

- Array of data type *Type* and length *Length*
- Contiguously allocated region of *Length* * **sizeof**(*Type*) bytes in memory

char string[12];



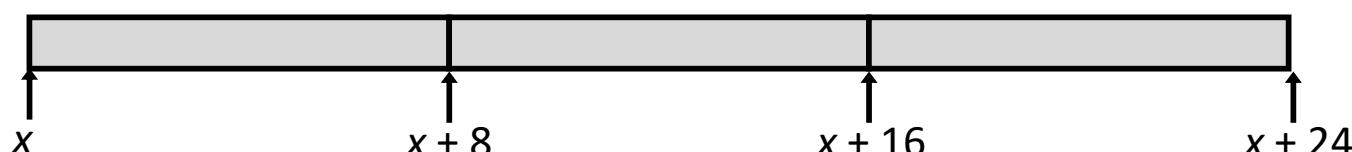
int val[5];



double a[3];



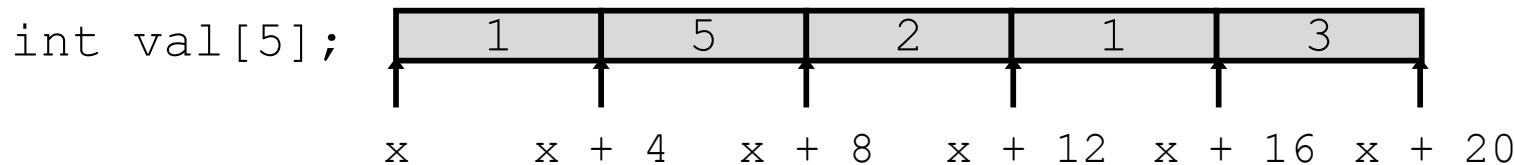
char *p[3];



Array Access

C declaration **Type name[Length]** ;

- Array of data type *Type* and length *Length*
- Identifier **name** acts like¹ a pointer to array element 0



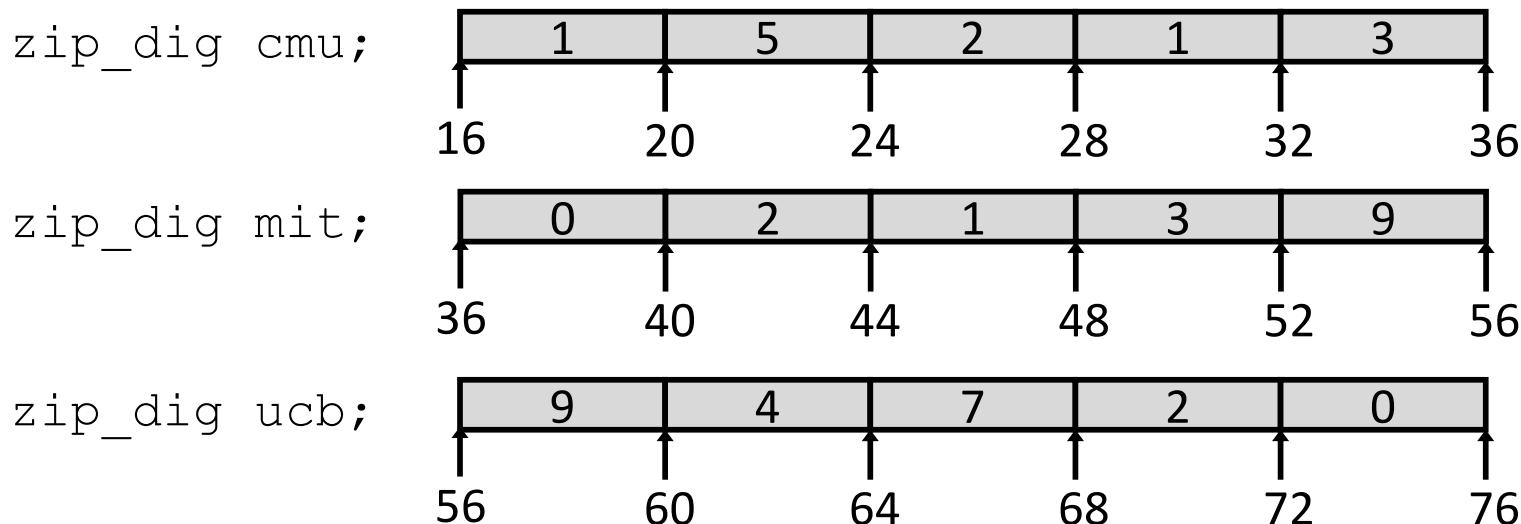
Expression	Type	Value	
<code>val[4]</code>	<code>int</code>	3	
<code>val[5]</code>	<code>int</code>	??	// access past end
<code>* (val+3)</code>	<code>int</code>	1	// same as <code>val[3]</code>
<code>val</code>	<code>int *</code>	<code>x</code>	
<code>val+1</code>	<code>int *</code>	<code>x + 4</code>	
<code>&val[2]</code>	<code>int *</code>	<code>x + 8</code>	// same as <code>val+2</code>
<code>val + i</code>	<code>int *</code>	<code>x + 4*i</code>	// same as <code>&val[i]</code>

¹ in most contexts (but not all)

Array Example

```
#define ZLEN 5
typedef int zip_dig[ZLEN];

zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```



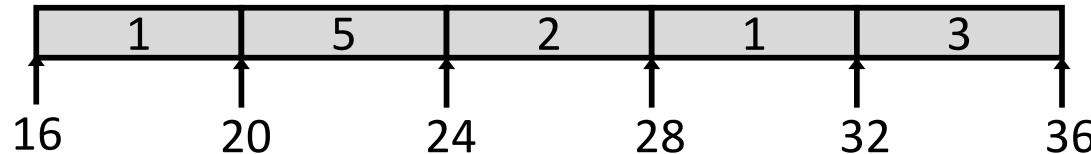
Declaration “`zip_dig cmu`” equivalent to “`int cmu[5]`”

Example arrays were allocated in successive 20 byte blocks

- Not guaranteed to happen in general

Array Accessing Example

```
zip_dig cmu;
```



```
int get_digit  
    (zip_dig z, int digit)  
{  
    return z[digit];  
}
```

x86-64

```
# %rdi = z  
# %rsi = digit  
movl (%rdi,%rsi,4), %eax # z[digit]
```

- Register **%rdi** contains starting address of array
- Register **%rsi** contains array index
- Desired digit at **%rdi + 4 * %rsi**
- Use memory reference **(%rdi,%rsi,4)**

Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl    $0, %eax  
jmp     .L3  
.L4:  
    addl    $1, (%rdi,%rax,4)  
    addq    $1, %rax  
.L3:  
    cmpq    $4, %rax  
    jbe     .L4  
rep; ret
```

Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl $0, %eax          # i = 0  
jmp .L3                # goto middle  
.L4:                  # loop:  
    addl $1, (%rdi,%rax,4) # z[i]++  
    addq $1, %rax          # i++  
.L3:                  # middle  
    cmpq $4, %rax          # i:4  
    jbe .L4                # if <=, goto loop  
rep; ret
```

Understanding Pointers & Arrays #1

Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
<code>int A1[3]</code>						
<code>int *A2</code>						

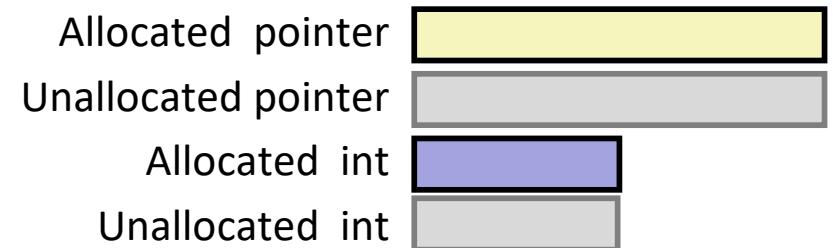
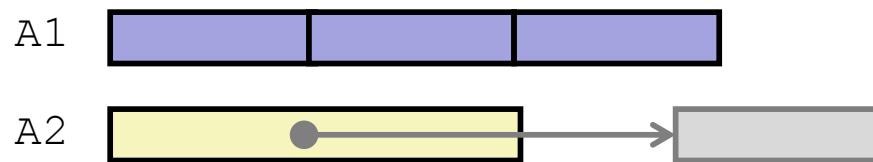
Comp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by sizeof

Understanding Pointers & Arrays #1

Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
int A1[3]						
int *A2						



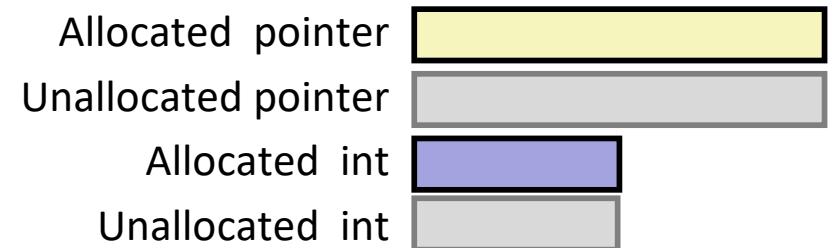
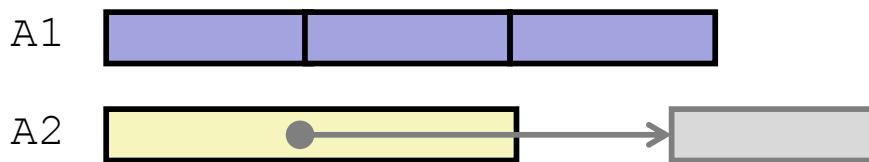
Comp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by sizeof

Understanding Pointers & Arrays #1

Decl	A1 , A2			*A1 , *A2		
	Comp	Bad	Size	Comp	Bad	Size
int A1[3]	Y	N	12	Y	N	4
int *A2	Y	N	8	Y	Y	4



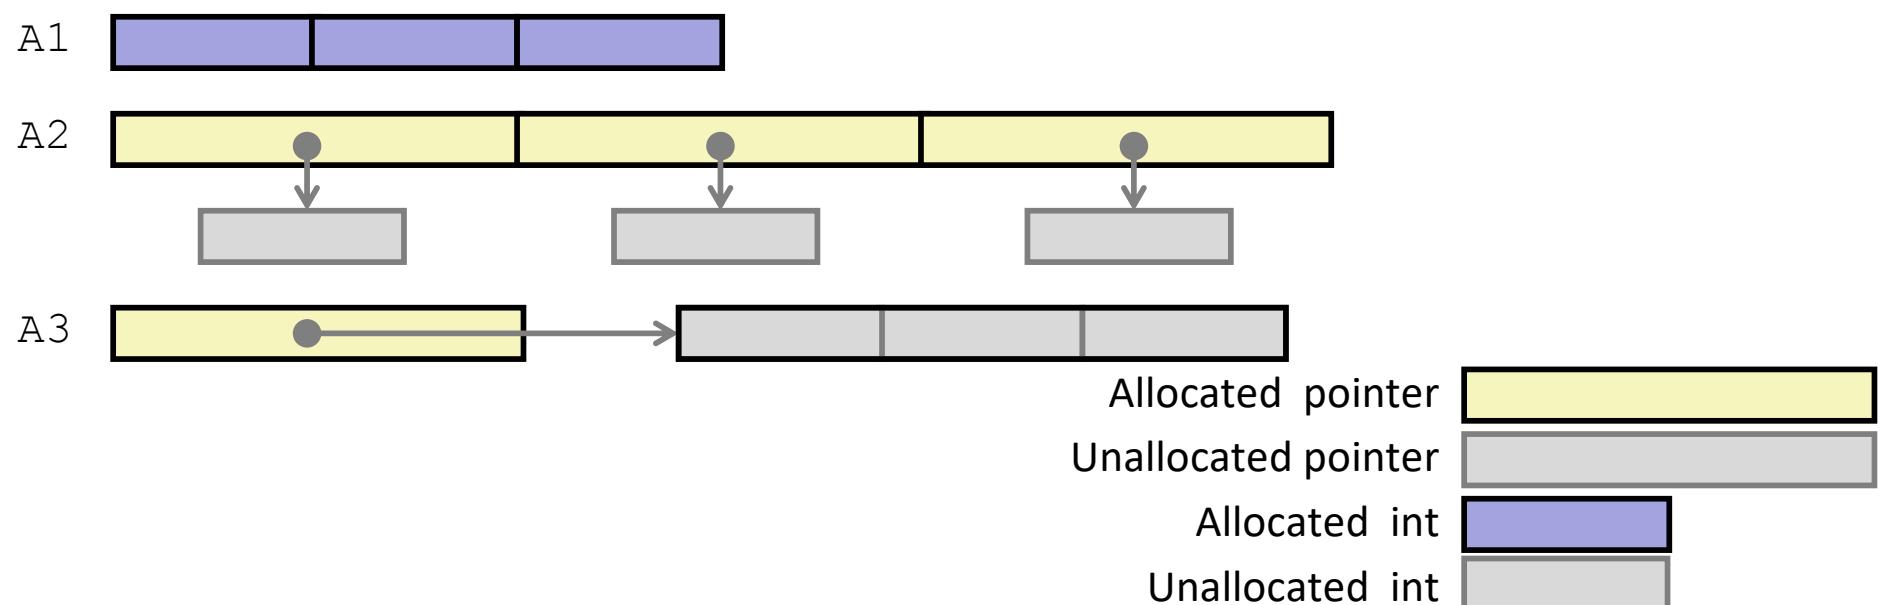
Comp: Compiles (Y/N)

Bad: Possible bad pointer reference (Y/N)

Size: Value returned by sizeof

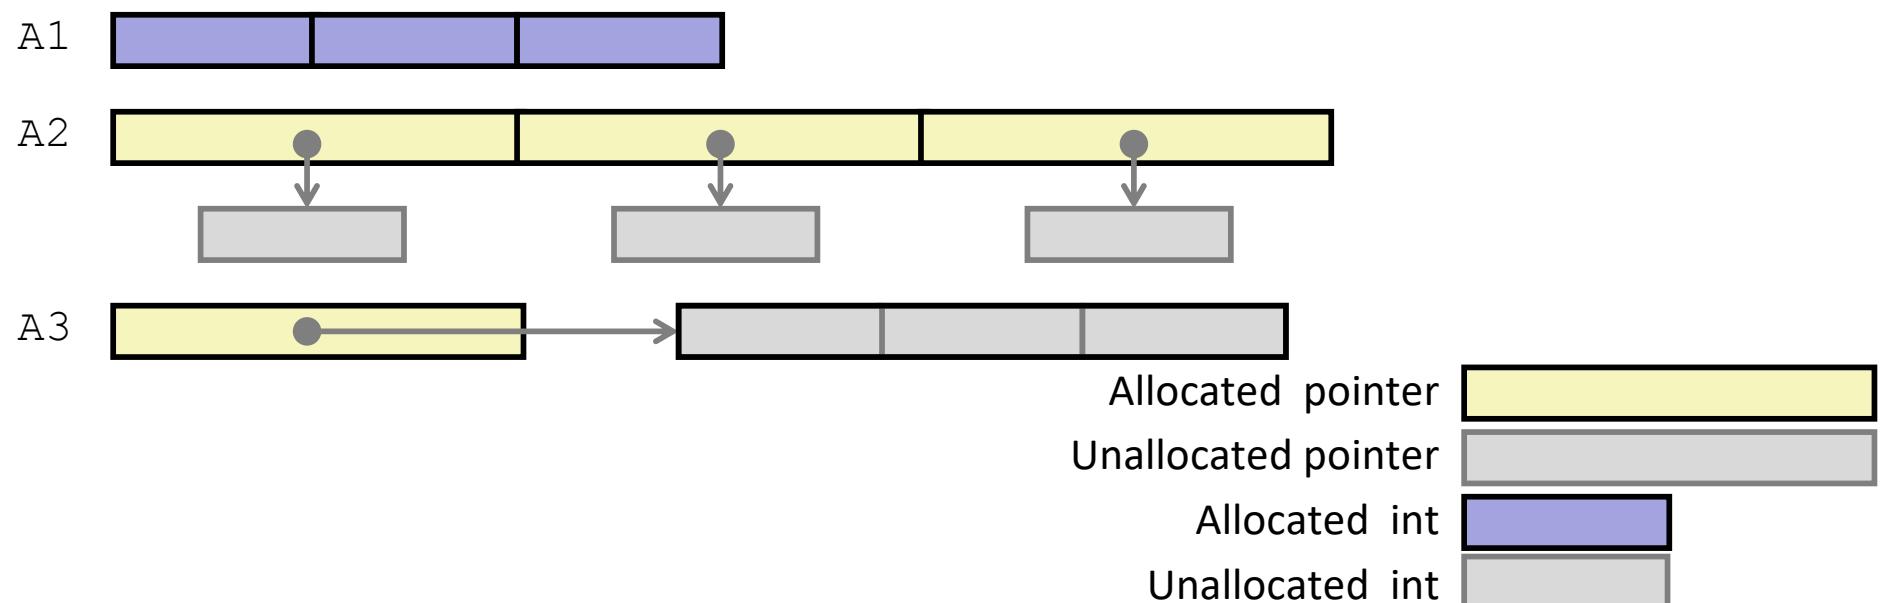
Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>									
<code>int *A2[3]</code>									
<code>int (*A3)[3]</code>									



Understanding Pointers & Arrays #2

Decl	<i>An</i>			<i>*An</i>			<i>**An</i>		
	Cmp	Bad	Size	Cmp	Bad	Size	Cmp	Bad	Size
<code>int A1[3]</code>	Y	N	12	Y	N	4	N	-	-
<code>int *A2[3]</code>	Y	N	24	Y	N	8	Y	Y	4
<code>int (*A3)[3]</code>	Y	N	8	Y	Y	12	Y	Y	4



Multidimensional (Nested) Arrays

Declaration

$T \ A[R][C];$

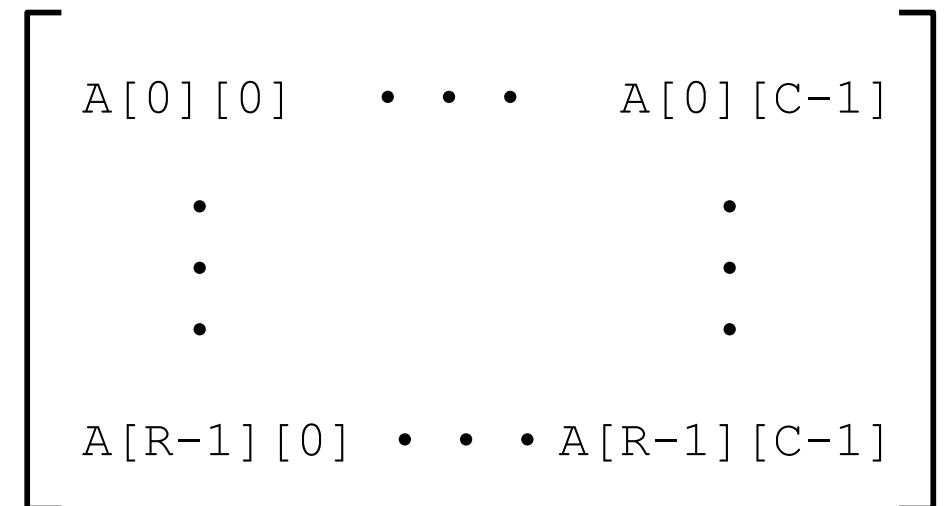
- 2D array of data type T
- R rows, C columns

Array Size

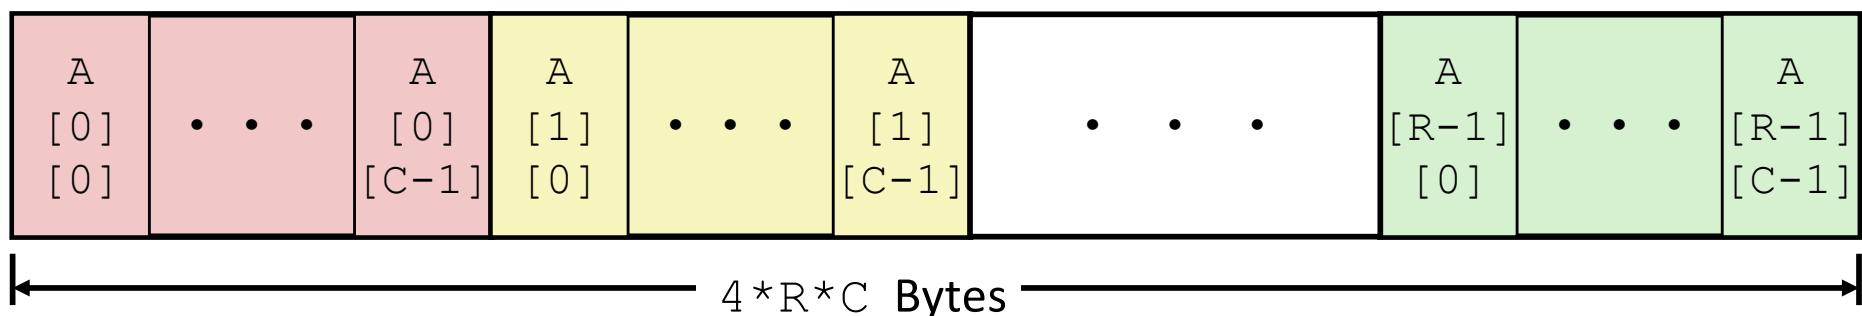
- $R * C * \text{sizeof}(T)$ bytes

Arrangement

- Row-Major Ordering



```
int A[R][C];
```

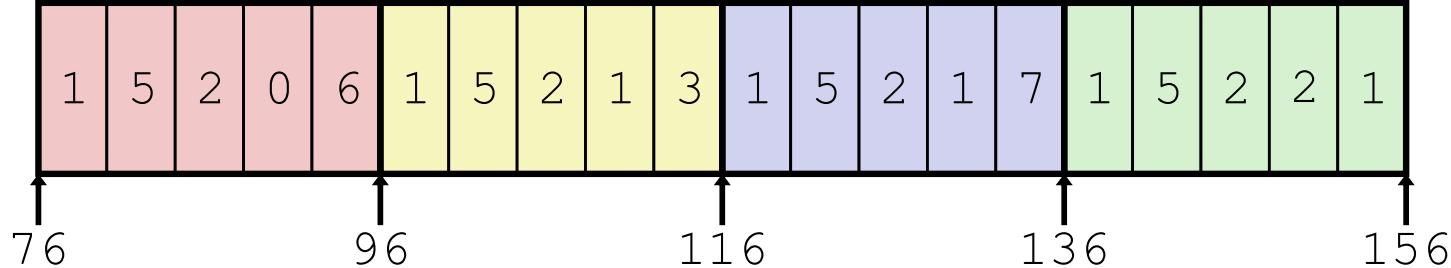


Nested Array Example

```
#define PCOUNT 4
typedef int zip_dig[5];

zip_dig pgh [ PCOUNT ] =
    {{1, 5, 2, 0, 6 },
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```

`zip_dig
pgh[4];`



“`zip_dig pgh[4]`” equivalent to “`int pgh[4][5]`”

- Variable `pgh`: array of 4 elements, allocated contiguously
- Each element is an array of 5 `int`'s, allocated contiguously

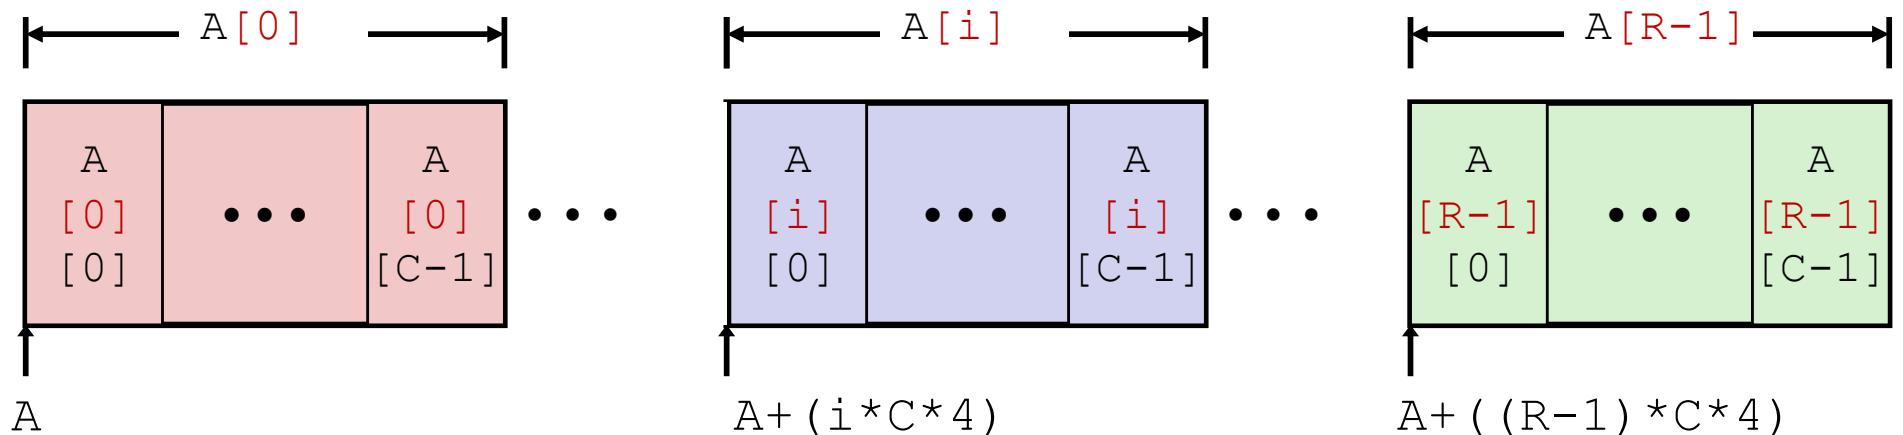
“Row-Major” ordering of all elements in memory

Nested Array Row Access

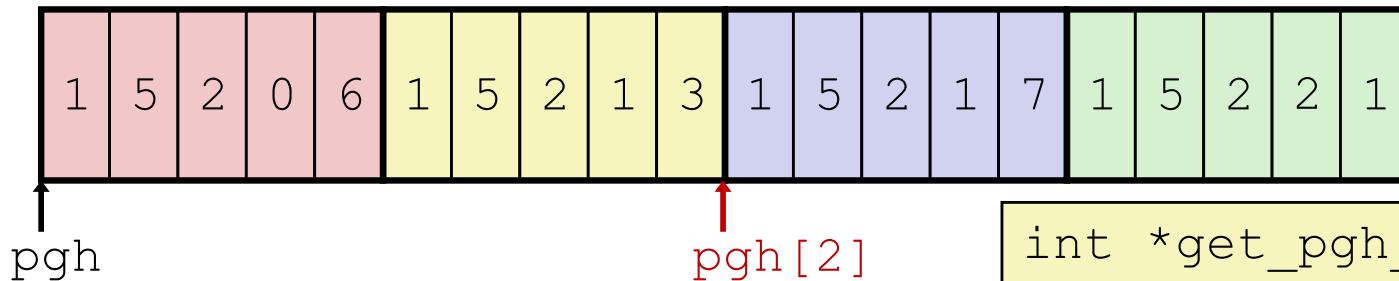
Row Vectors

- $\mathbf{A[i]}$ is array of C elements of type T
- Starting address $\mathbf{A} + i * (C * \text{sizeof}(T))$

```
int A[R][C];
```



Nested Array Row Access Code



```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
# %rdi = index
leaq (%rdi,%rdi,4),%rax # 5 * index
leaq pgh(,%rax,4),%rax # pgh + (20 * index)
```

Row Vector

- **pgh[index]** is array of 5 **int's**
- Starting address **pgh+20*index**

Machine Code

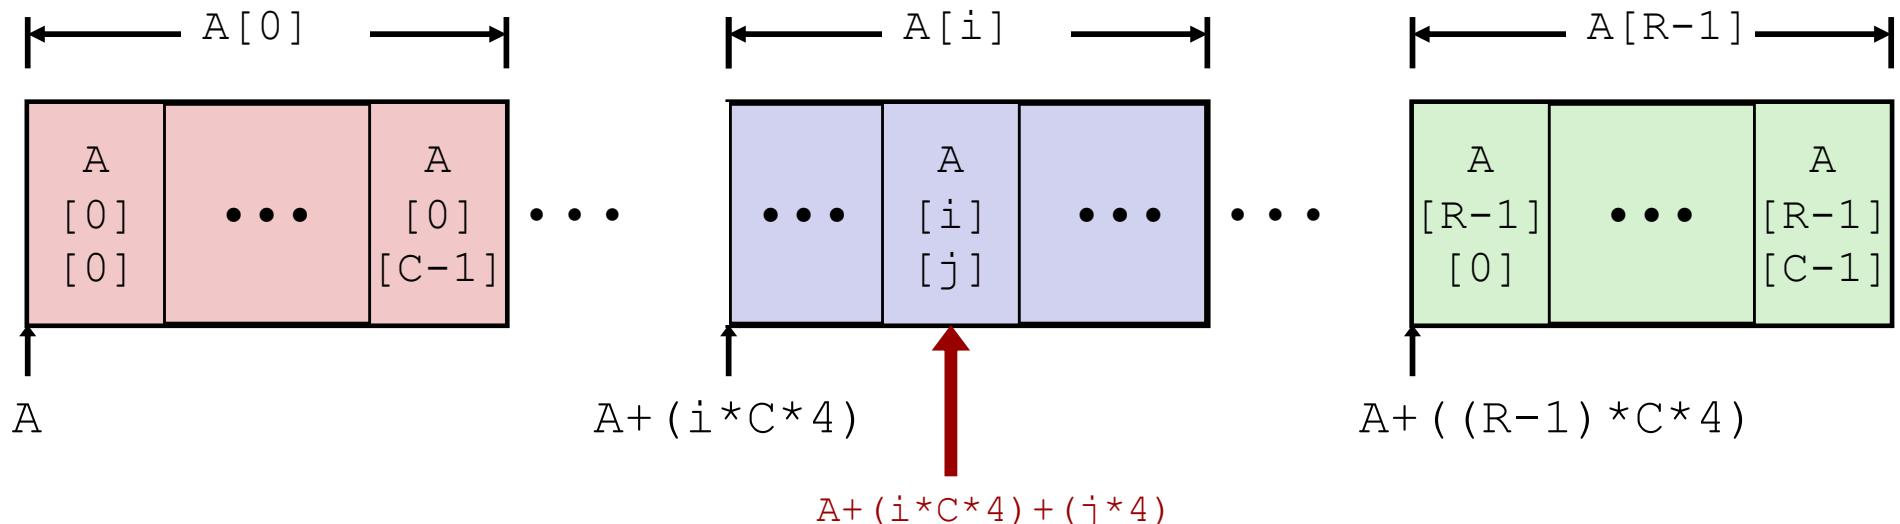
- Computes and returns address
- Compute as **pgh + 4*(index+4*index)**

Nested Array Element Access

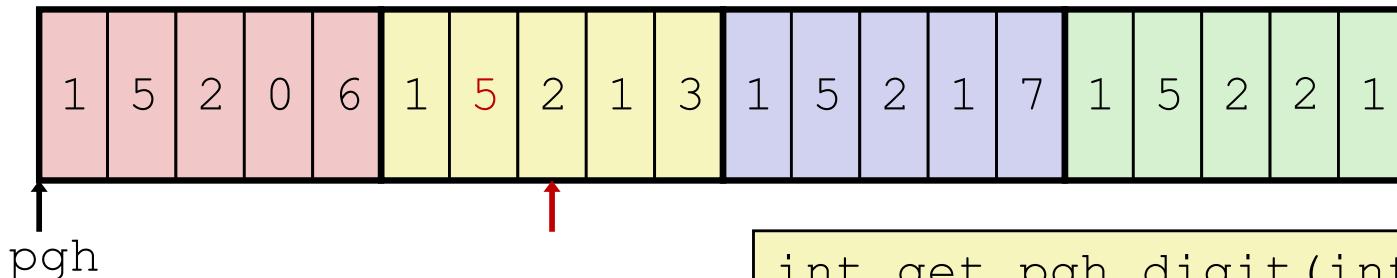
Array Elements

- $A[i][j]$ is element of type T , which requires K bytes
- Address $A + i * (C * K) + j * K$
 $= A + (i * C + j) * K$

```
int A[R][C];
```



Nested Array Element Access Code



```
int get_pgh_digit(int index, int dig)
{
    return pgh[index][dig];
}
```

```
leaq (%rdi,%rdi,4), %rax      # 5*index
addl %rax, %rsi                # 5*index+dig
movl pgh(,%rsi,4), %eax       # M[pgh + 4*(5*index+dig)]
```

Array Elements

- **pgh[index][dig]** is **int**
- Address: **pgh + 20*index + 4*dig**
 $= \text{pgh} + 4 * (5 * \text{index} + \text{dig})$

Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

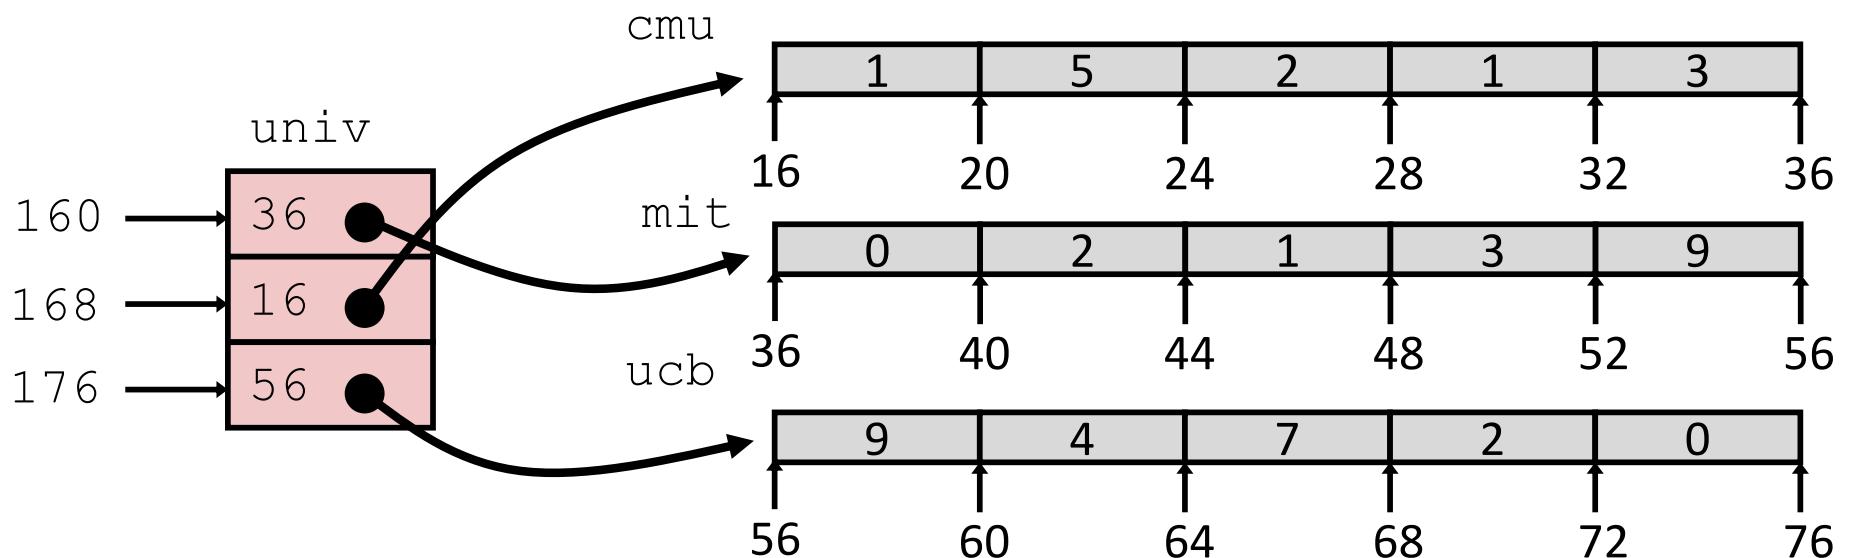
```
#define UCOUNT 3
int *univ[UCOUNT] = {mit, cmu, ucb};
```

Variable `univ` denotes array of 3 elements

Each element is a pointer

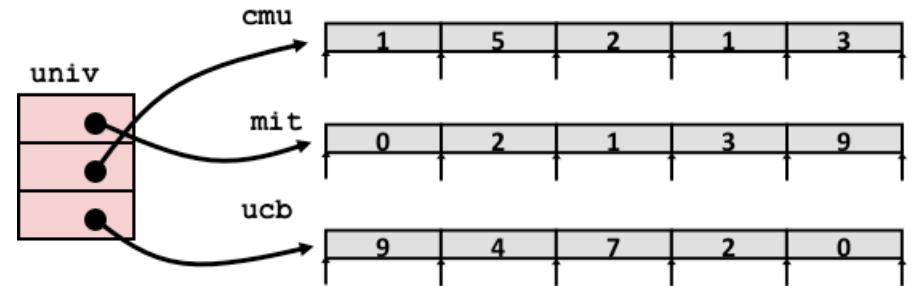
- 8 bytes

Each pointer points to array of `int`'s



Element Access in Multi-Level Array

```
int get_univ_digit
    (size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq    univ(%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
ret
```

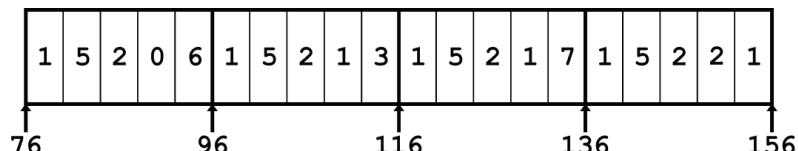
Computation

- Element access **Mem[Mem[univ+8*index]+4*digit]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array

Array Element Accesses

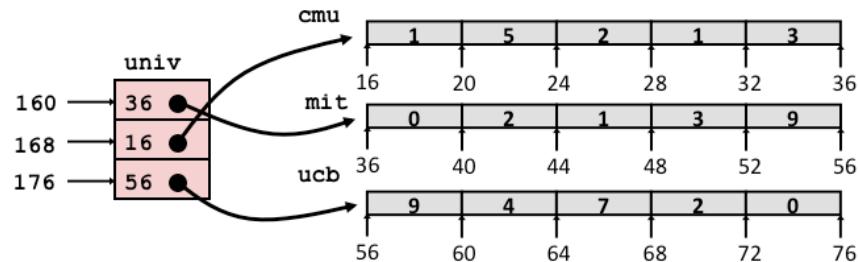
Nested array

```
int get_pgh_digit
    (size_t index, size_t digit)
{
    return pgh[index][digit];
}
```



Multi-level array

```
int get_univ_digit
    (size_t index, size_t digit)
{
    return univ[index][digit];
}
```



Accesses looks similar in C, but address computations very different:

`Mem[pgh+20*index+4*digit]`

`Mem[Mem[univ+8*index]+4*digit]`

$N \times N$ Matrix Code

Fixed dimensions

- Know value of N at compile time

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element A[i][j] */
int fix_ele(fix_matrix A,
            size_t i, size_t j)
{
    return A[i][j];
}
```

Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element A[i][j] */
int vec_ele(size_t n, int *A,
            size_t i, size_t j)
{
    return A[IDX(n,i,j)];
}
```

Variable dimensions, implicit indexing

- Not in K&R; added to language in 1999

```
/* Get element A[i][j] */
int var_ele(size_t n, int A[n][n],
            size_t i, size_t j) {
    return A[i][j];
}
```

16 X 16 Matrix Access

■ Array Elements

- `int A[16][16];`
- Address `A + i * (C * K) + j * K`
- $C = 16, K = 4$

```
/* Get element A[i][j] */  
int fix_ele(fix_matrix A, size_t i, size_t j) {  
    return A[i][j];  
}
```

```
# A in %rdi, i in %rsi, j in %rdx  
salq    $6, %rsi           # 64*i  
addq    %rsi, %rdi         # A + 64*i  
movl    (%rdi,%rdx,4), %eax # Mem[A + 64*i + 4*j]  
ret
```

$n \times n$ Matrix Access

■ Array Elements

- `size_t n;`
- `int A[n][n];`
- Address $A + i * (C * K) + j * K$
- $C = n, K = 4$
- Must perform integer multiplication

```
/* Get element A[i][j] */
int var_ele(size_t n, int A[n][n], size_t i, size_t j)
{
    return A[i][j];
}
```

```
# n in %rdi, A in %rsi, i in %rdx, j in %rcx
imulq    %rdx, %rdi          # n*i
leaq     (%rsi,%rdi,4), %rax # A + 4*n*i
movl     (%rax,%rcx,4), %eax # Mem[A + 4*n*i + 4*j]
ret
```

Example: Array Access

```
#include <stdio.h>

#define ZLEN 5
#define PCOUNT 4

typedef int zip_dig[ZLEN];

int main(int argc, char** argv) {
    zip_dig pgm[PCOUNT] =
        {{1, 5, 2, 0, 6},
         {1, 5, 2, 1, 3},
         {1, 5, 2, 1, 7},
         {1, 5, 2, 2, 1}};

    int *linear_zip = (int *) pgm;
    int *zip2 = (int *) pgm[2];
    int result =
        pgm[0][0] +
        linear_zip[7] +
        *(linear_zip + 8) +
        zip2[1];
    printf("result: %d\n", result);
    return 0;
}
```

```
linux> ./array
```

Example: Array Access

```
#include <stdio.h>

#define ZLEN 5
#define PCOUNT 4

typedef int zip_dig[ZLEN];

int main(int argc, char** argv) {
    zip_dig pgm[PCOUNT] =
        {{1, 5, 2, 0, 6},
         {1, 5, 2, 1, 3},
         {1, 5, 2, 1, 7},
         {1, 5, 2, 2, 1}};

    int *linear_zip = (int *) pgm;
    int *zip2 = (int *) pgm[2];
    int result =
        pgm[0][0] +
        linear_zip[7] +
        *(linear_zip + 8) +
        zip2[1];
    printf("result: %d\n", result);
    return 0;
}
```

```
linux> ./array
result: 9
```

Today

Procedures

- Stack Structure
- Calling Conventions
 - Passing control
 - Passing data
 - Managing local data
- Illustration of Recursion

Arrays

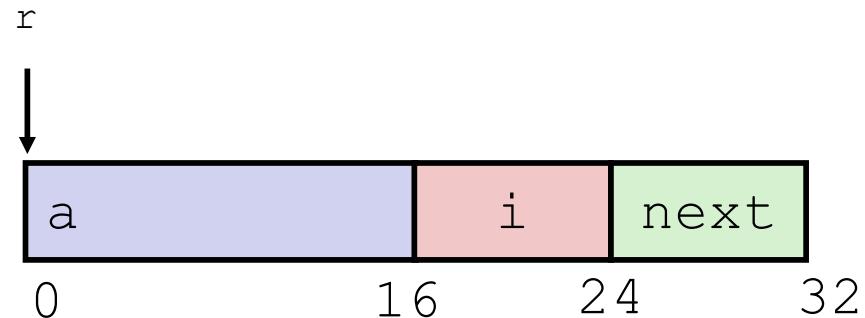
- One-dimensional
- Multi-dimensional (nested)
- Multi-level

Structures

- Allocation
- Access
- Alignment

Structure Representation

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



Structure represented as block of memory

- Big enough to hold all the fields

Fields ordered according to declaration

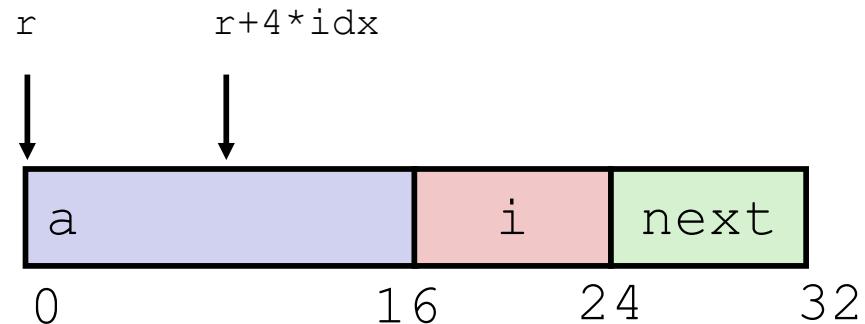
- Even if another ordering could be more compact

Compiler determines overall size + positions of fields

- In assembly, we see only offsets, not field names

Generating Pointer to Structure Member

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Compute as `r + 4*idx`

```
int *get_ap
(struct rec *r, size_t idx)
{
    return &r->a[idx];
}
```

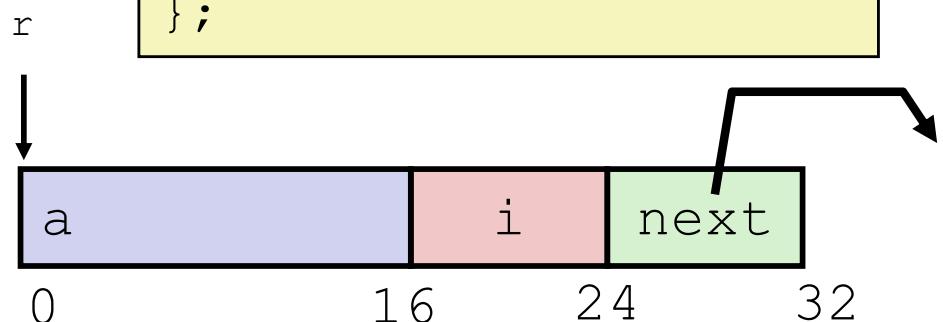
```
# r in %rdi, idx in %rsi
leaq (%rdi,%rsi,4), %rax
ret
```

Following Linked List #1

C Code

```
long length(struct rec*r) {
    long len = 0L;
    while (r) {
        len++;
        r = r->next;
    }
    return len;
}
```

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```



Register	Value
<code>%rdi</code>	<code>r</code>
<code>%rax</code>	<code>len</code>

Loop assembly code

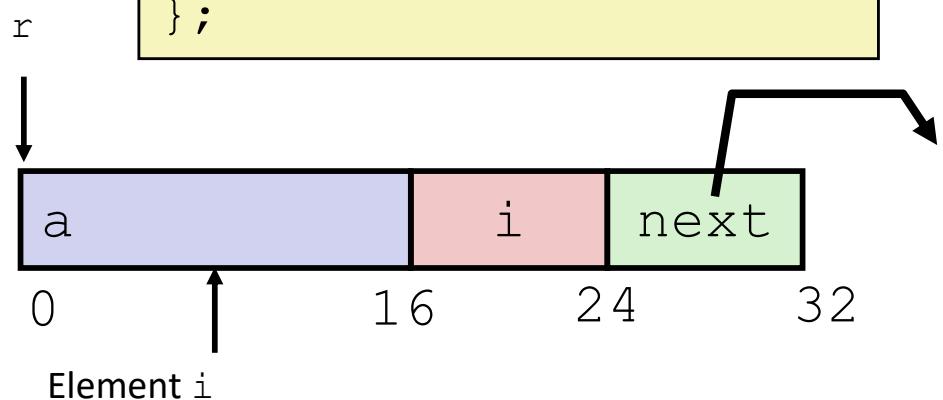
```
.L11:                                # loop:
    addq    $1, %rax                 # len ++
    movq    24(%rdi), %rdi          # r = Mem[r+24]
    testq   %rdi, %rdi              # Test r
    jne     .L11                   # If != 0, goto loop
```

Following Linked List #2

C Code

```
void set_val
    (struct rec *r, int val)
{
    while (r) {
        size_t i = r->i;
        // No bounds check
        r->a[r->i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    size_t i;
    struct rec *next;
};
```

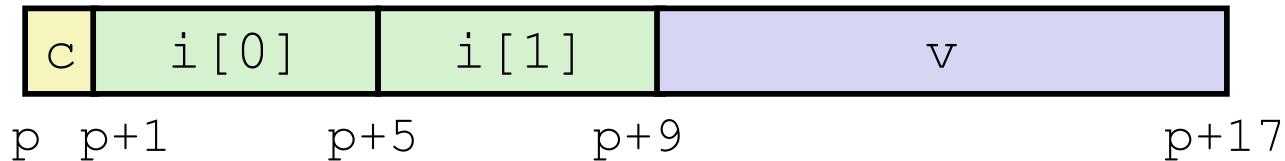


Register	Value
%rdi	<code>r</code>
%rsi	<code>val</code>

```
.L11:                                # loop:
    movq 16(%rdi), %rax             #     i = Mem[r+16]
    movl %esi, (%rdi,%rax,4)       #     Mem[r+4*i] = val
    movq 24(%rdi), %rdi            #     r = Mem[r+24]
    testq %rdi, %rdi               #     Test r
    jne   .L11                      #     if !=0 goto loop
```

Structures & Alignment

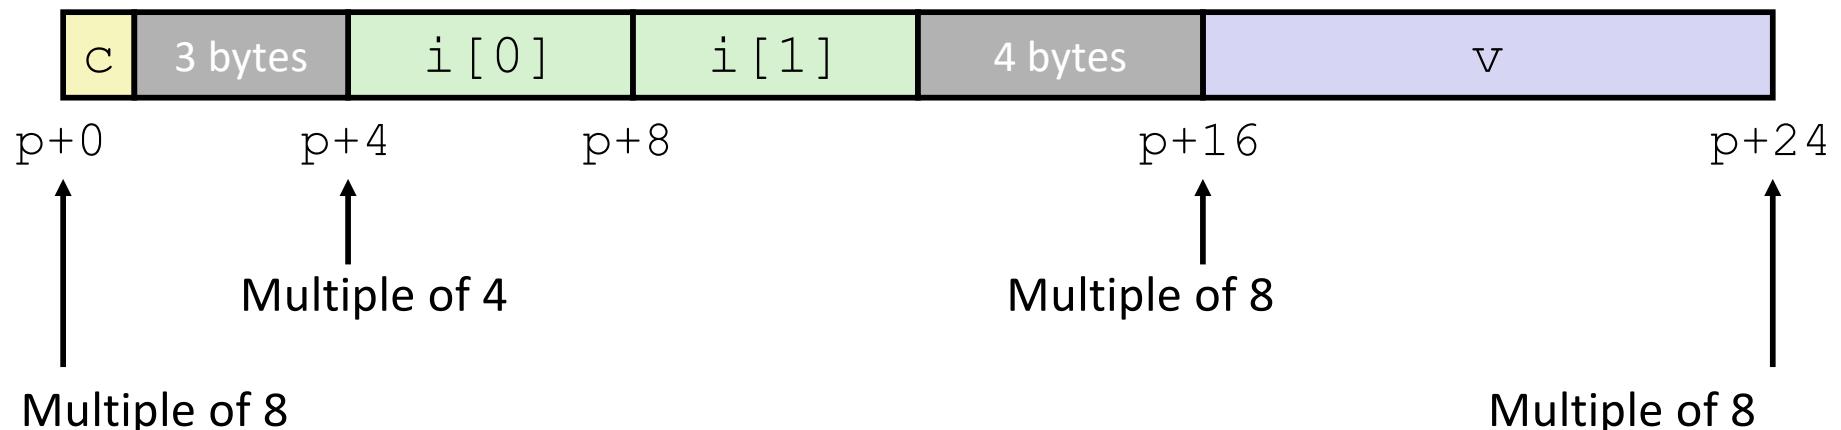
Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Aligned Data

- Primitive data type requires B bytes implies
Address must be multiple of B



Alignment Principles

Aligned Data

- Primitive data type requires B bytes
- Address must be multiple of B
- Required on some machines; advised on x86-64

Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans cache lines (64 bytes). Intel states should avoid crossing 16 byte boundaries.

[Cache lines will be discussed in Lecture 10.]

- Virtual memory trickier when datum spans 2 pages (4 KB pages)

[Virtual memory pages will be discussed in Lecture 17.]

Compiler

- Inserts gaps in structure to ensure correct alignment of fields

Specific Cases of Alignment (x86-64)

1 byte: `char`, ...

- no restrictions on address

2 bytes: `short`, ...

- lowest 1 bit of address must be 0_2

4 bytes: `int`, `float`, ...

- lowest 2 bits of address must be 00_2

8 bytes: `double`, `long`, `char *`, ...

- lowest 3 bits of address must be 000_2

Satisfying Alignment with Structures

Within structure:

- Must satisfy each element's alignment requirement

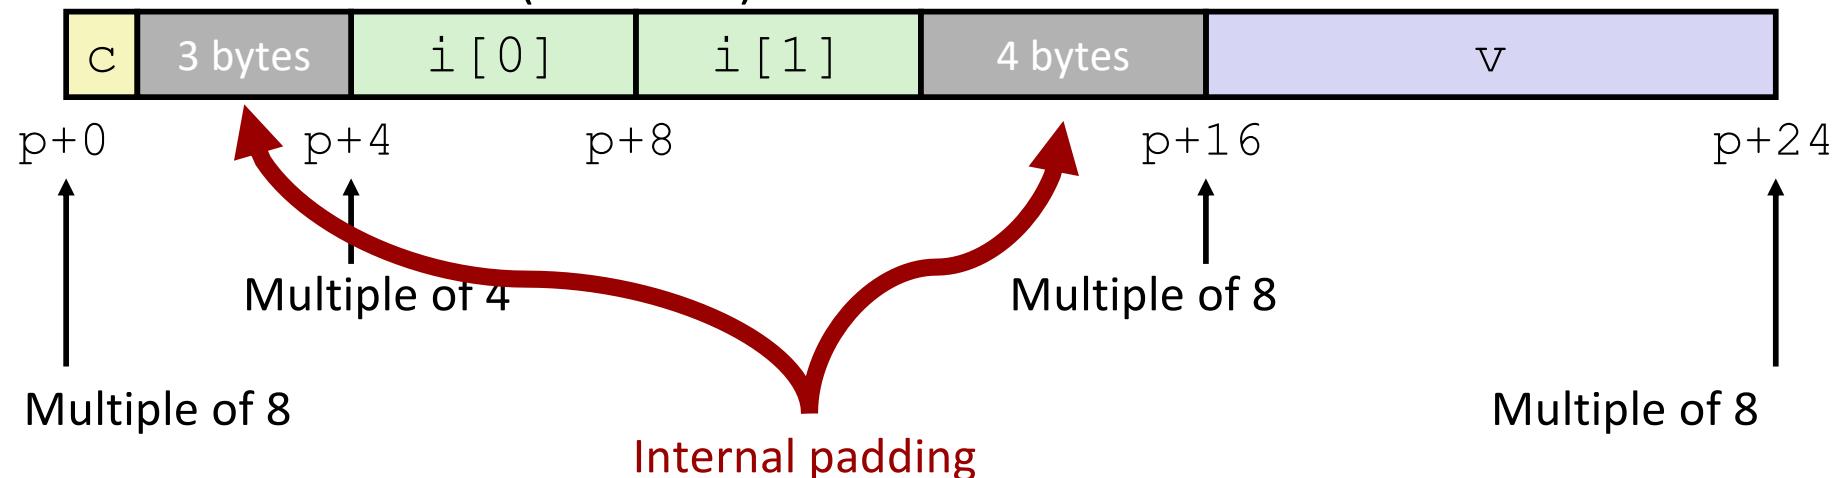
Overall structure placement

- Each structure has alignment requirement K
 - K = Largest alignment of any element
- Initial address & structure length must be multiples of K

Example:

- K = 8, due to **double** element

NOTE: K < sizeof(struct S1)



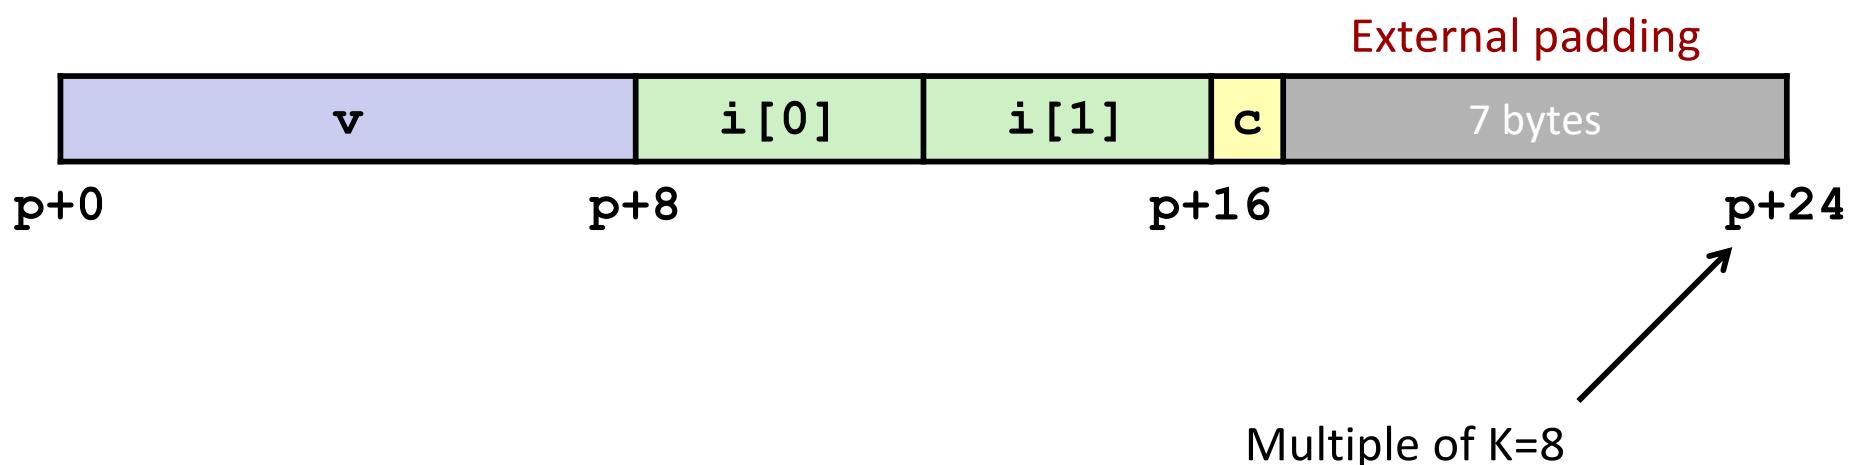
```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

Meeting Overall Alignment Requirement

For largest alignment requirement K

Overall structure must be multiple of K

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



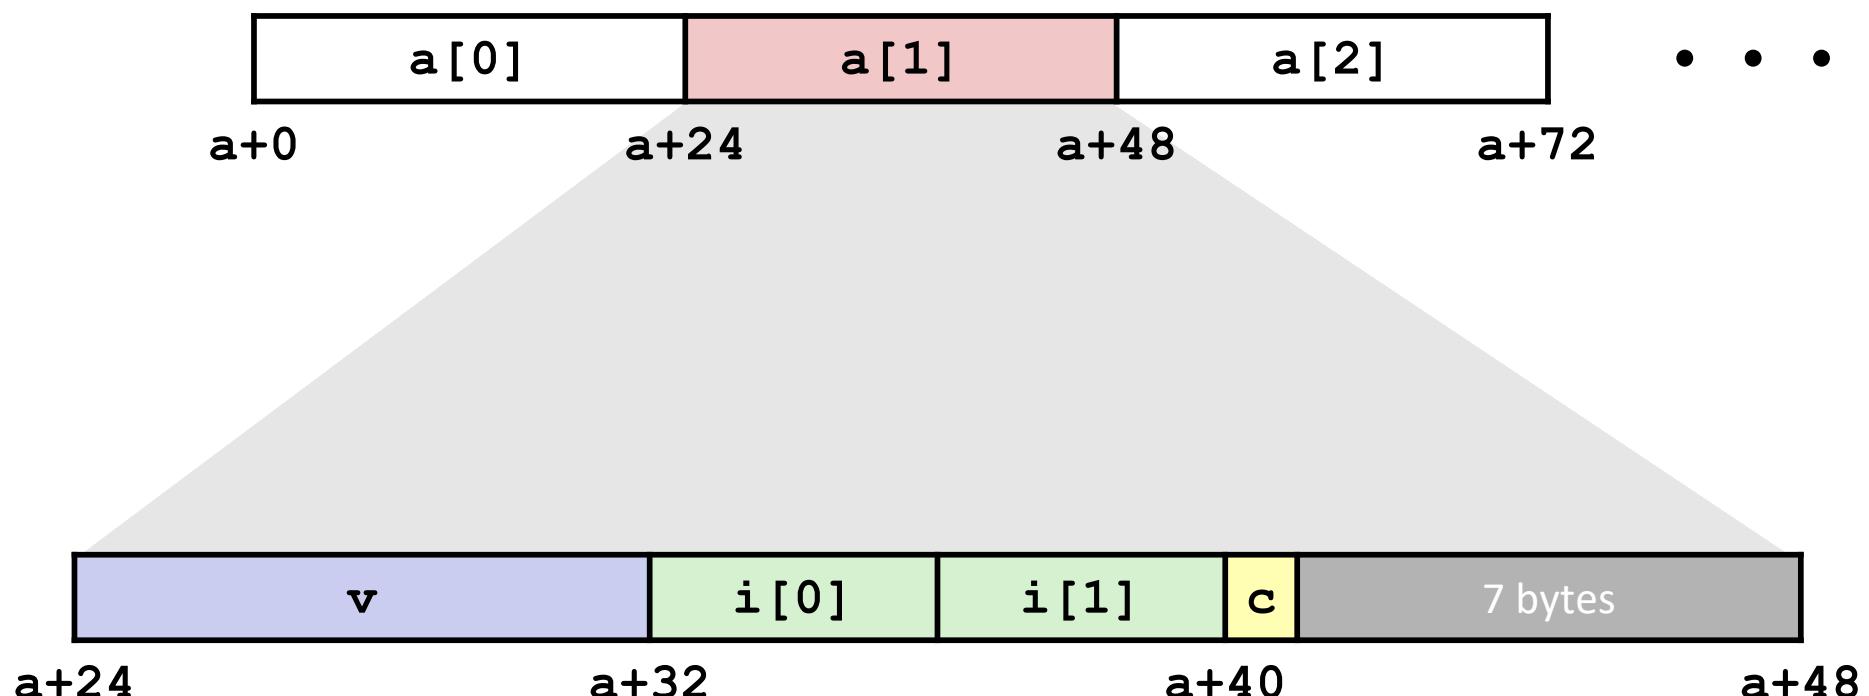
Arrays of Structures

No padding in between array elements

Overall structure length multiple of K

**Satisfy alignment requirement
for every element**

```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Array Elements

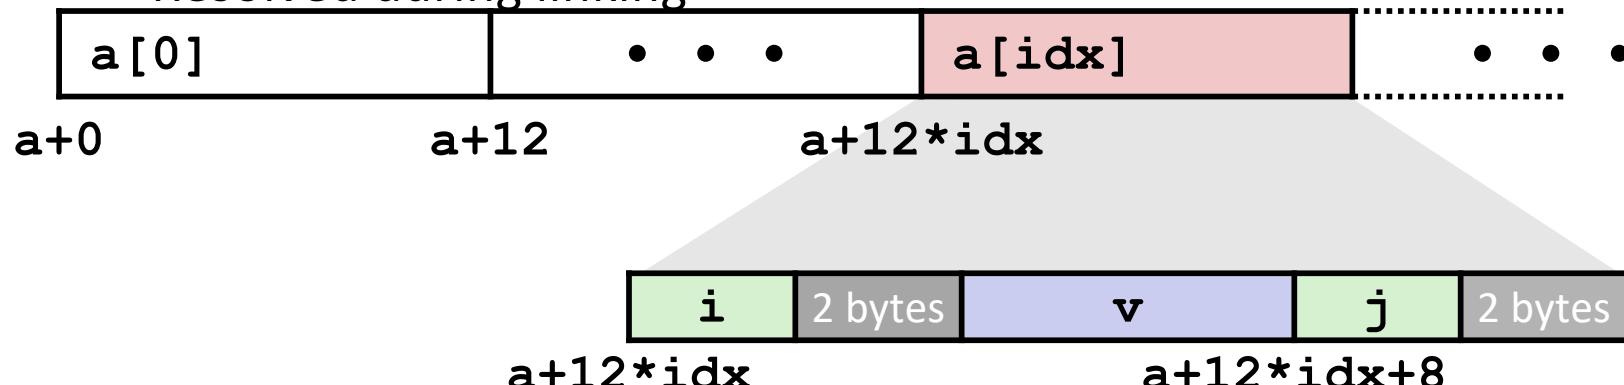
Compute array offset $12 * \text{idx}$

- `sizeof(S3)`, including alignment spacers

Element j is at offset 8 within structure

Assembler gives offset a+8

- Resolved during linking



```
short get_j(int idx)
{
    return a[idx].j;
}
```

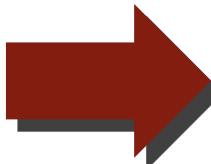
```
# %rdi = idx
leaq (%rdi,%rdi,2),%rax # 3*idx
movzwl a+8(%rax,4),%eax
```

```
struct S3 {
    short i;
    float v;
    short j;
} a[10];
```

Saving Space

Put large data types first

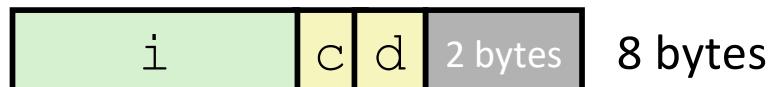
```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
char d;  
} *p;
```



Effect (largest alignment requirement K=4)



Summary

Arrays

- Elements packed into contiguous region of memory
- Use index arithmetic to locate individual elements

Structures

- Elements packed into single region of memory
- Access using offsets determined by compiler
- Possible require internal and external padding to ensure alignment

Combinations

- Can nest structure and array code arbitrarily

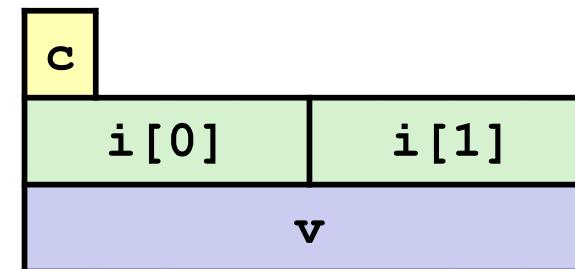
Union Allocation

Allocate according to largest element

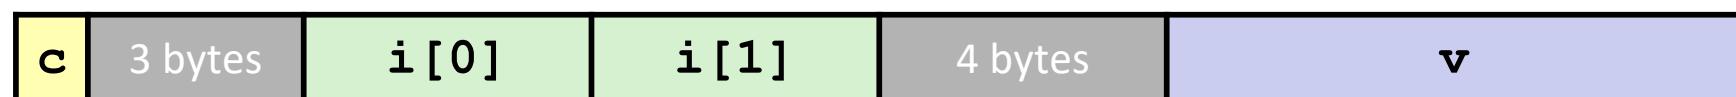
Can only use one field at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



up+0 up+4 up+8



sp+0 sp+4 sp+8 sp+16 sp+20 sp+24

Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as (**float**) **u** ?

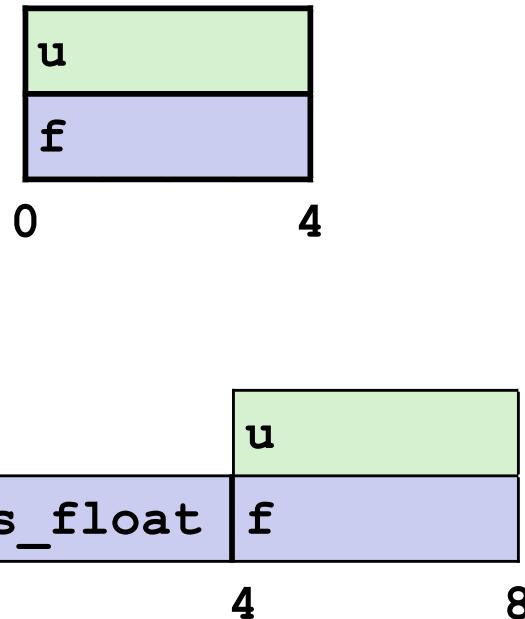
```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as (**unsigned**) **f** ?

Using Unions as Sum Types

```
typedef union {
    float f;
    unsigned u;
} num_t;

typedef struct {
    bool is_float;
    num_t val;
} value_t;
```



(technically `is_float` only takes 1 byte and then there's 3 bytes of padding)